



Shared memory machines; OpenMP

EECS 221: Intro to High-Performance Computing

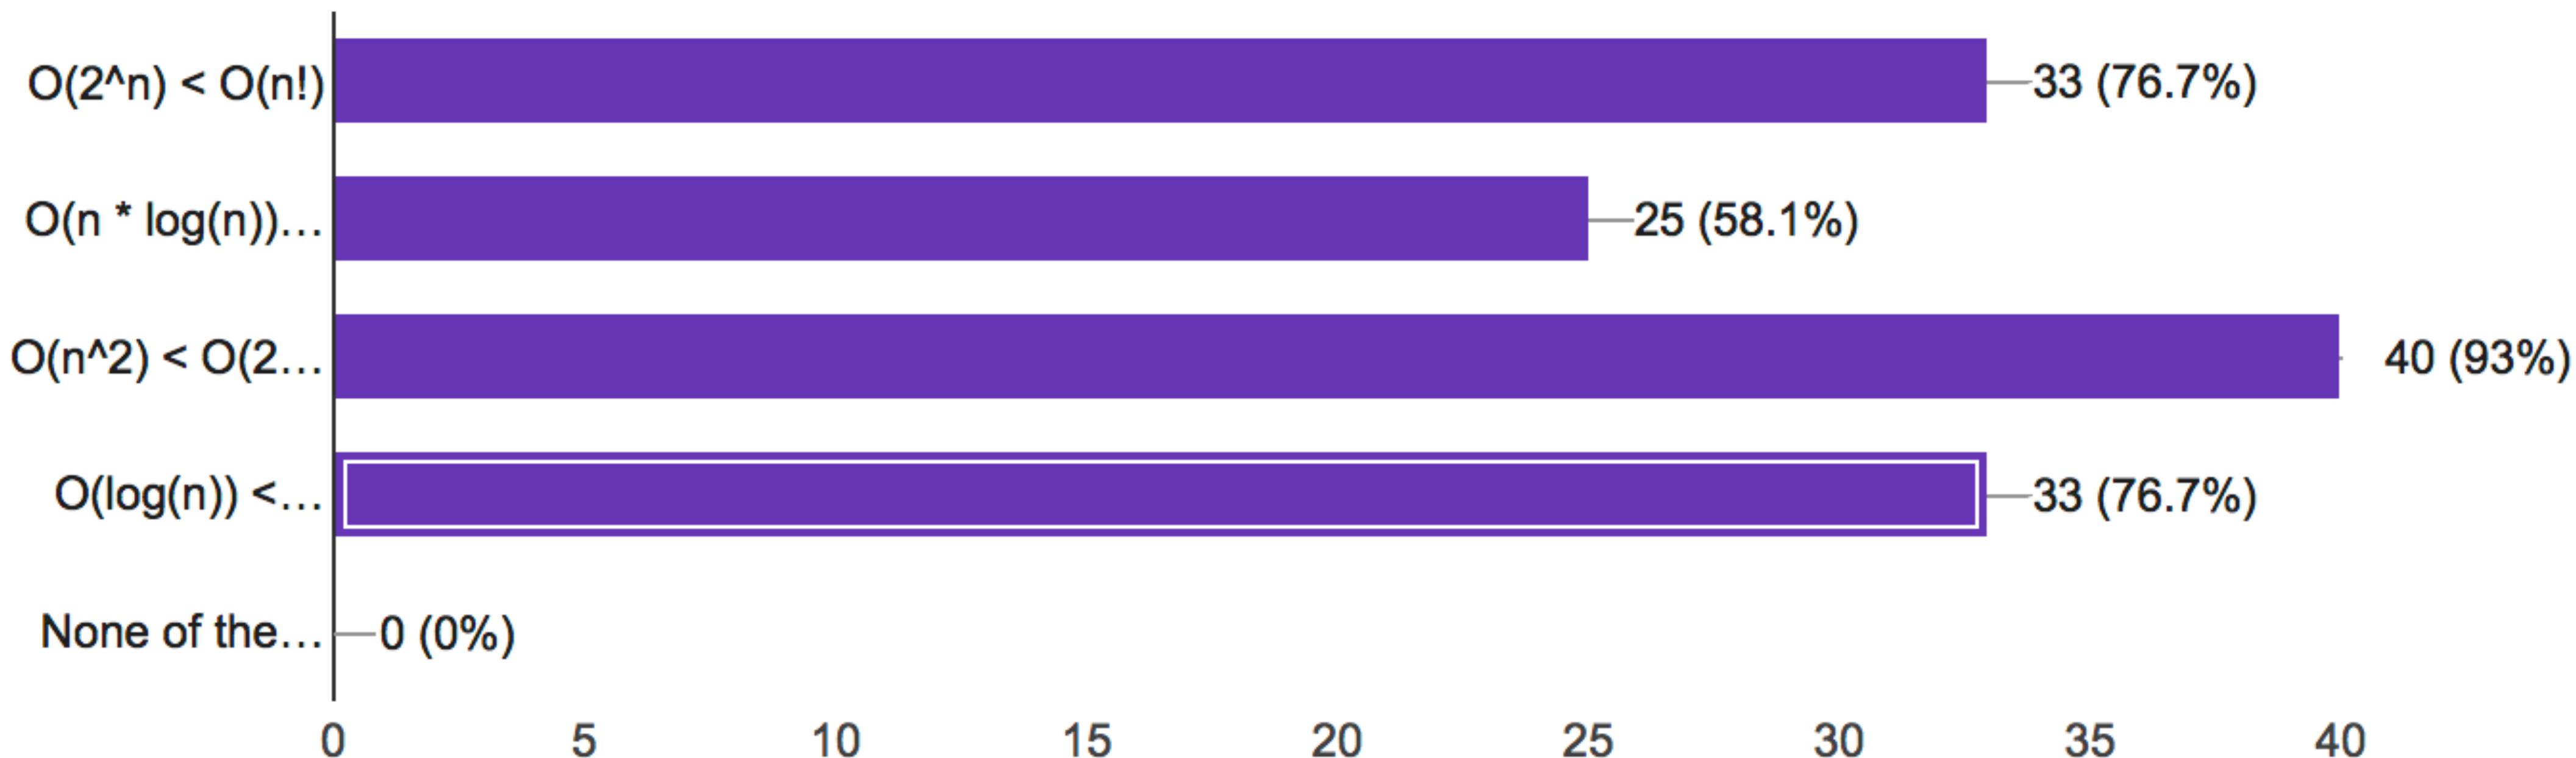
Aparna Chandramowlishwaran
April 13, 2017

Observations on Calibration Quiz/Survey

Observations

- ▶ Background in parallel computing: ~1/5 have some experience
- ▶ Programming language preference: C
- ▶ Basic Coding: Good!
- ▶ Big-O: Mostly OK, with **one common mistake**
- ▶ Performance bottleneck: Lots of answers

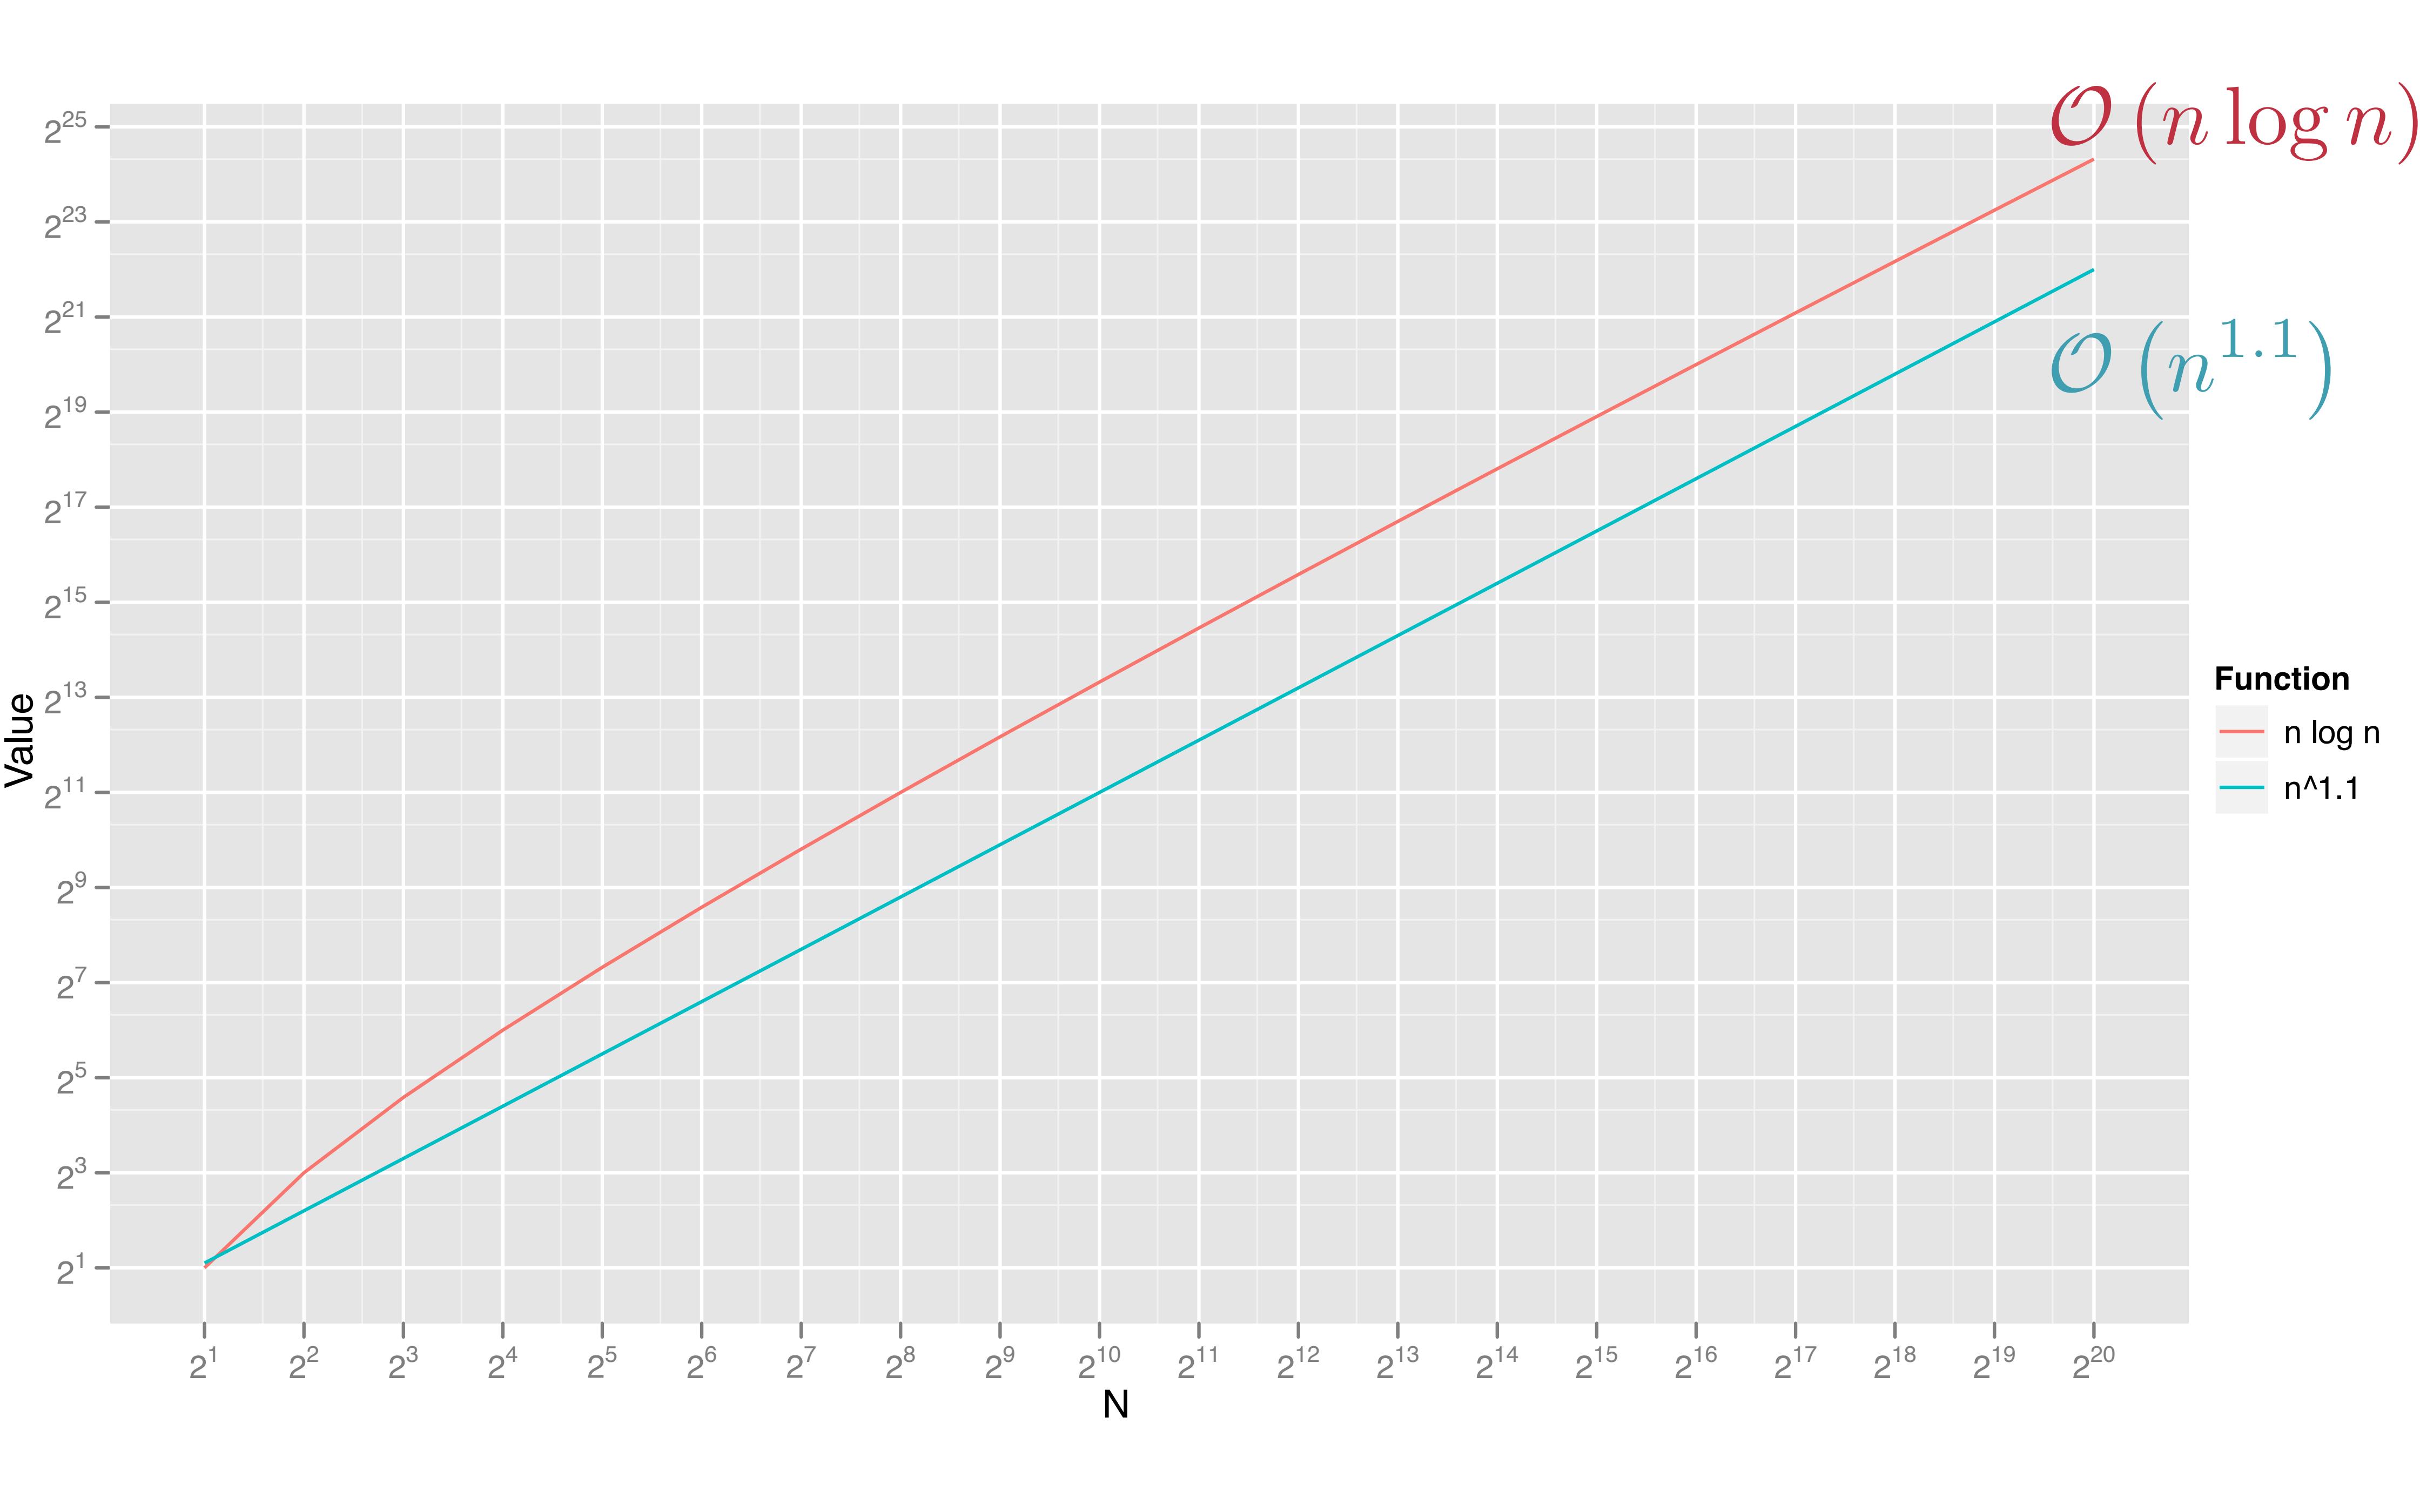
Which of these big-Oh statements is true? (43 responses)

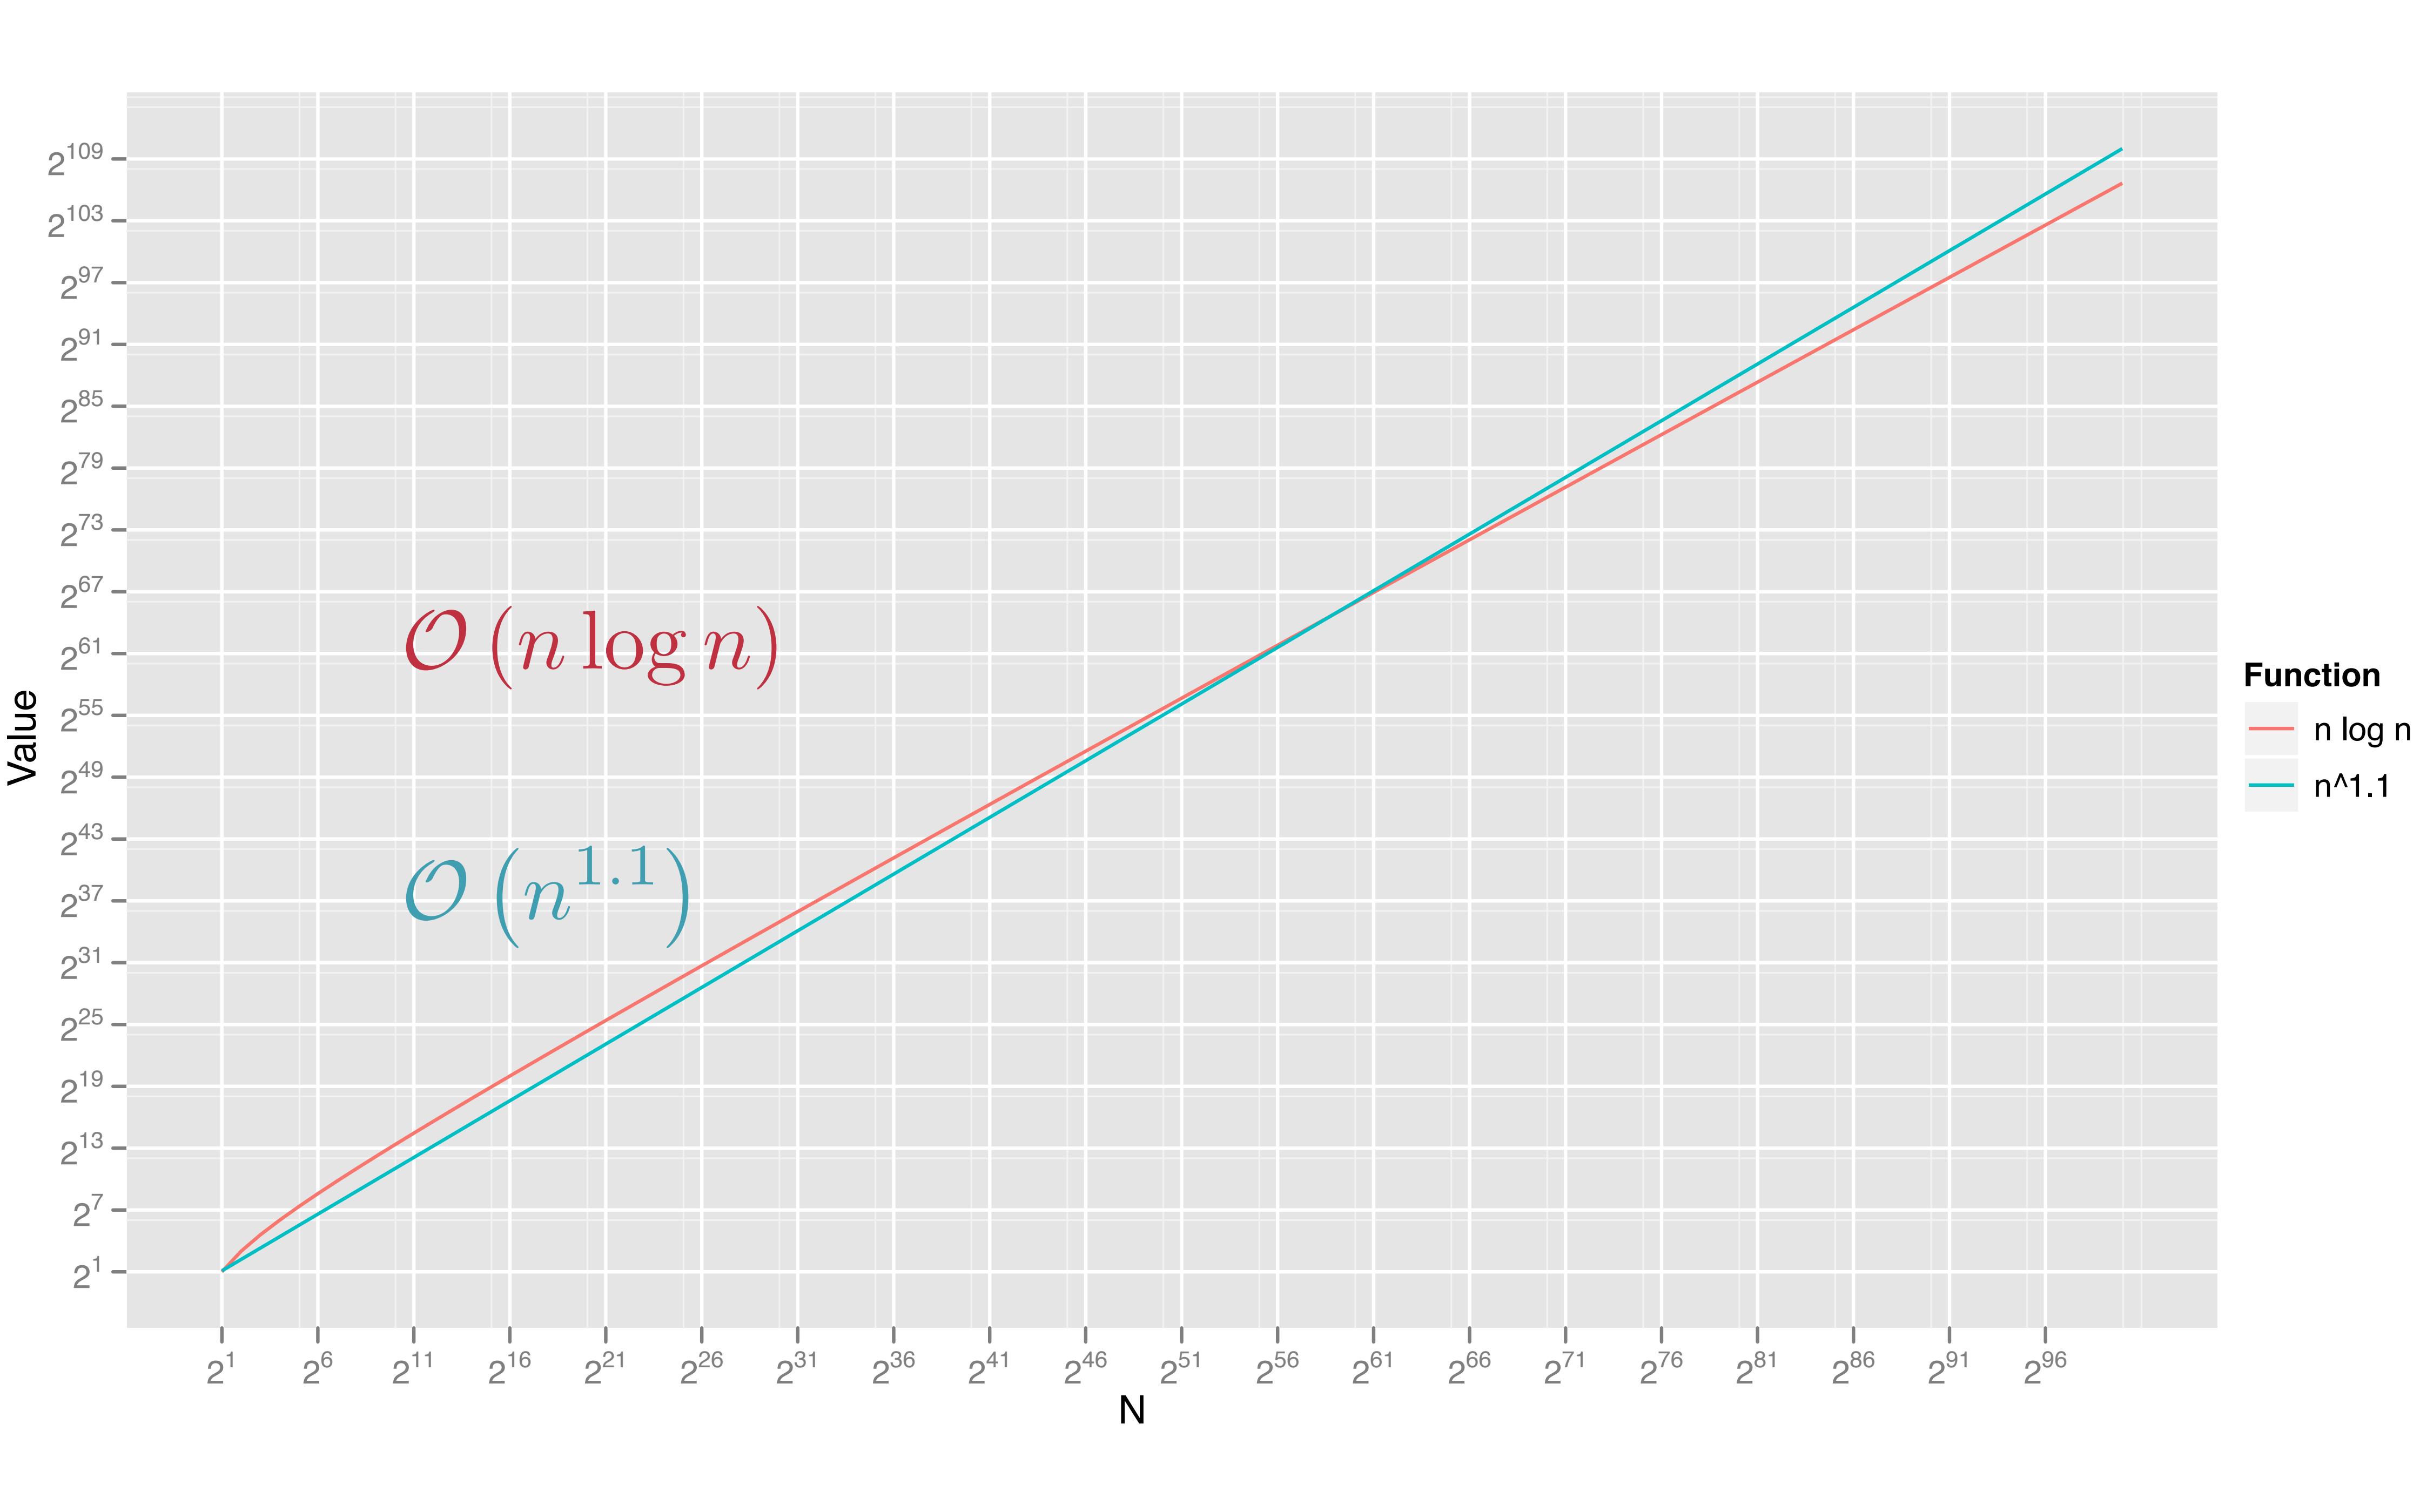


Which is larger?

$$\mathcal{O}(n \log n)$$

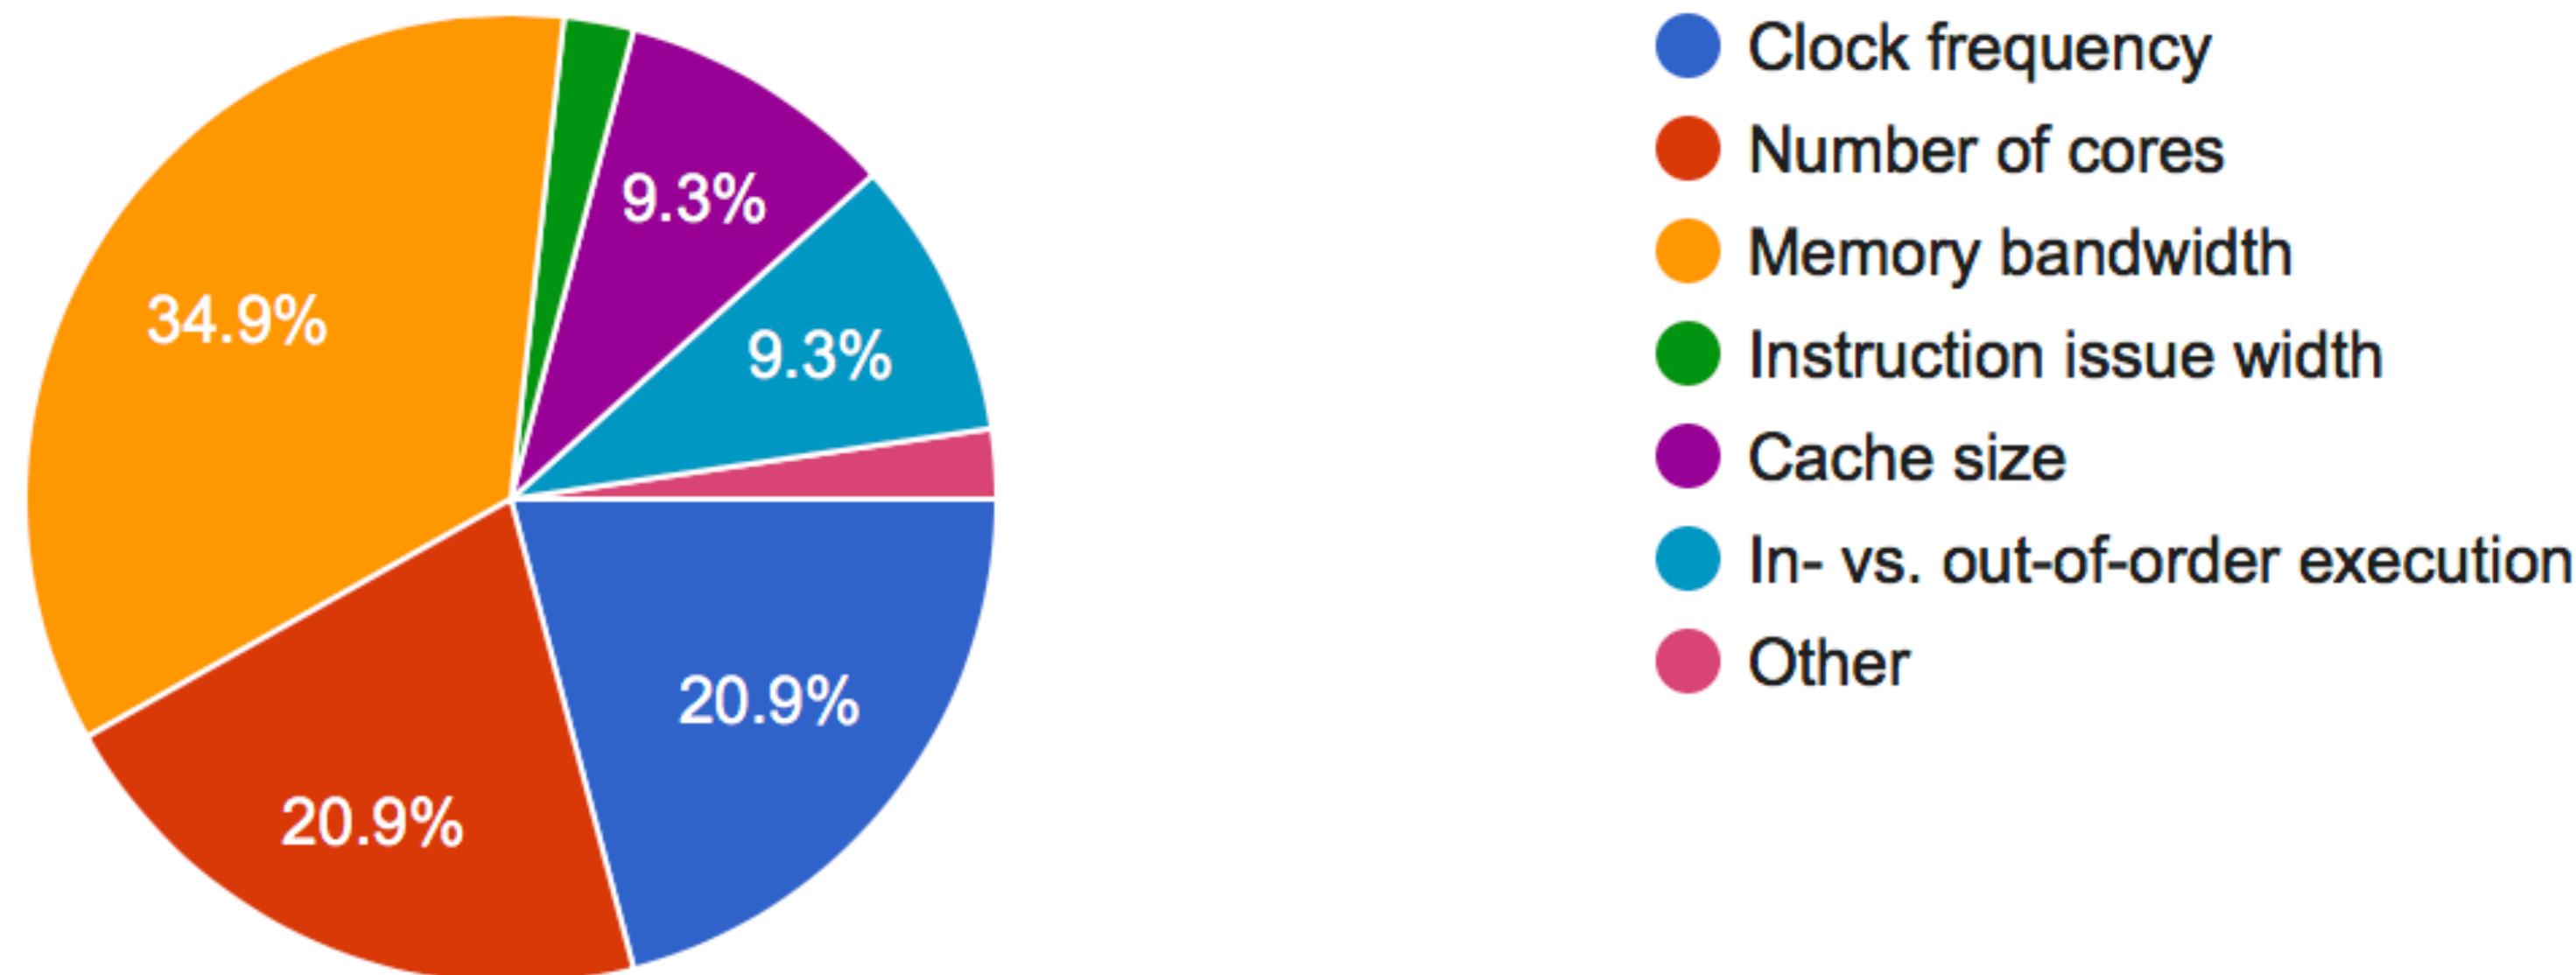
$$\mathcal{O}(n^{1.1})$$





In your vector-add code above, which of the following system characteristics is ***most likely*** to limit performance, when 'n' is very large?

(43 responses)



- ▶ Amdahl's Law:
Estimating best-case speedup
- ▶ Speedup vs efficiency:
Strong vs weak scaling (Gustafson's Law)
- ▶ *Shared memory programming model*
OpenMP

Amdahl's Law

- ▶ f = Fraction of work done **sequentially**
- ▶ $1-f$ = Fraction of work **parallelizable**
- ▶ p = Number of processors

$$S_p(f) = \frac{1}{f + \frac{1-f}{p}}$$

$$\lim_{p \rightarrow \infty} S_p(f) = \frac{1}{f}$$

Amdahl's Law

Exercise: Suppose sorting takes **70%** of the execution time of a program. Replace sequential sort with merge sort that **scales perfectly** on p processors. How many processors do you need to get a speedup of **4x**?

$$S_p(f) = \frac{1}{f + \frac{1-f}{p}}$$

$$\lim_{p \rightarrow \infty} S_p(f) = \frac{1}{f}$$

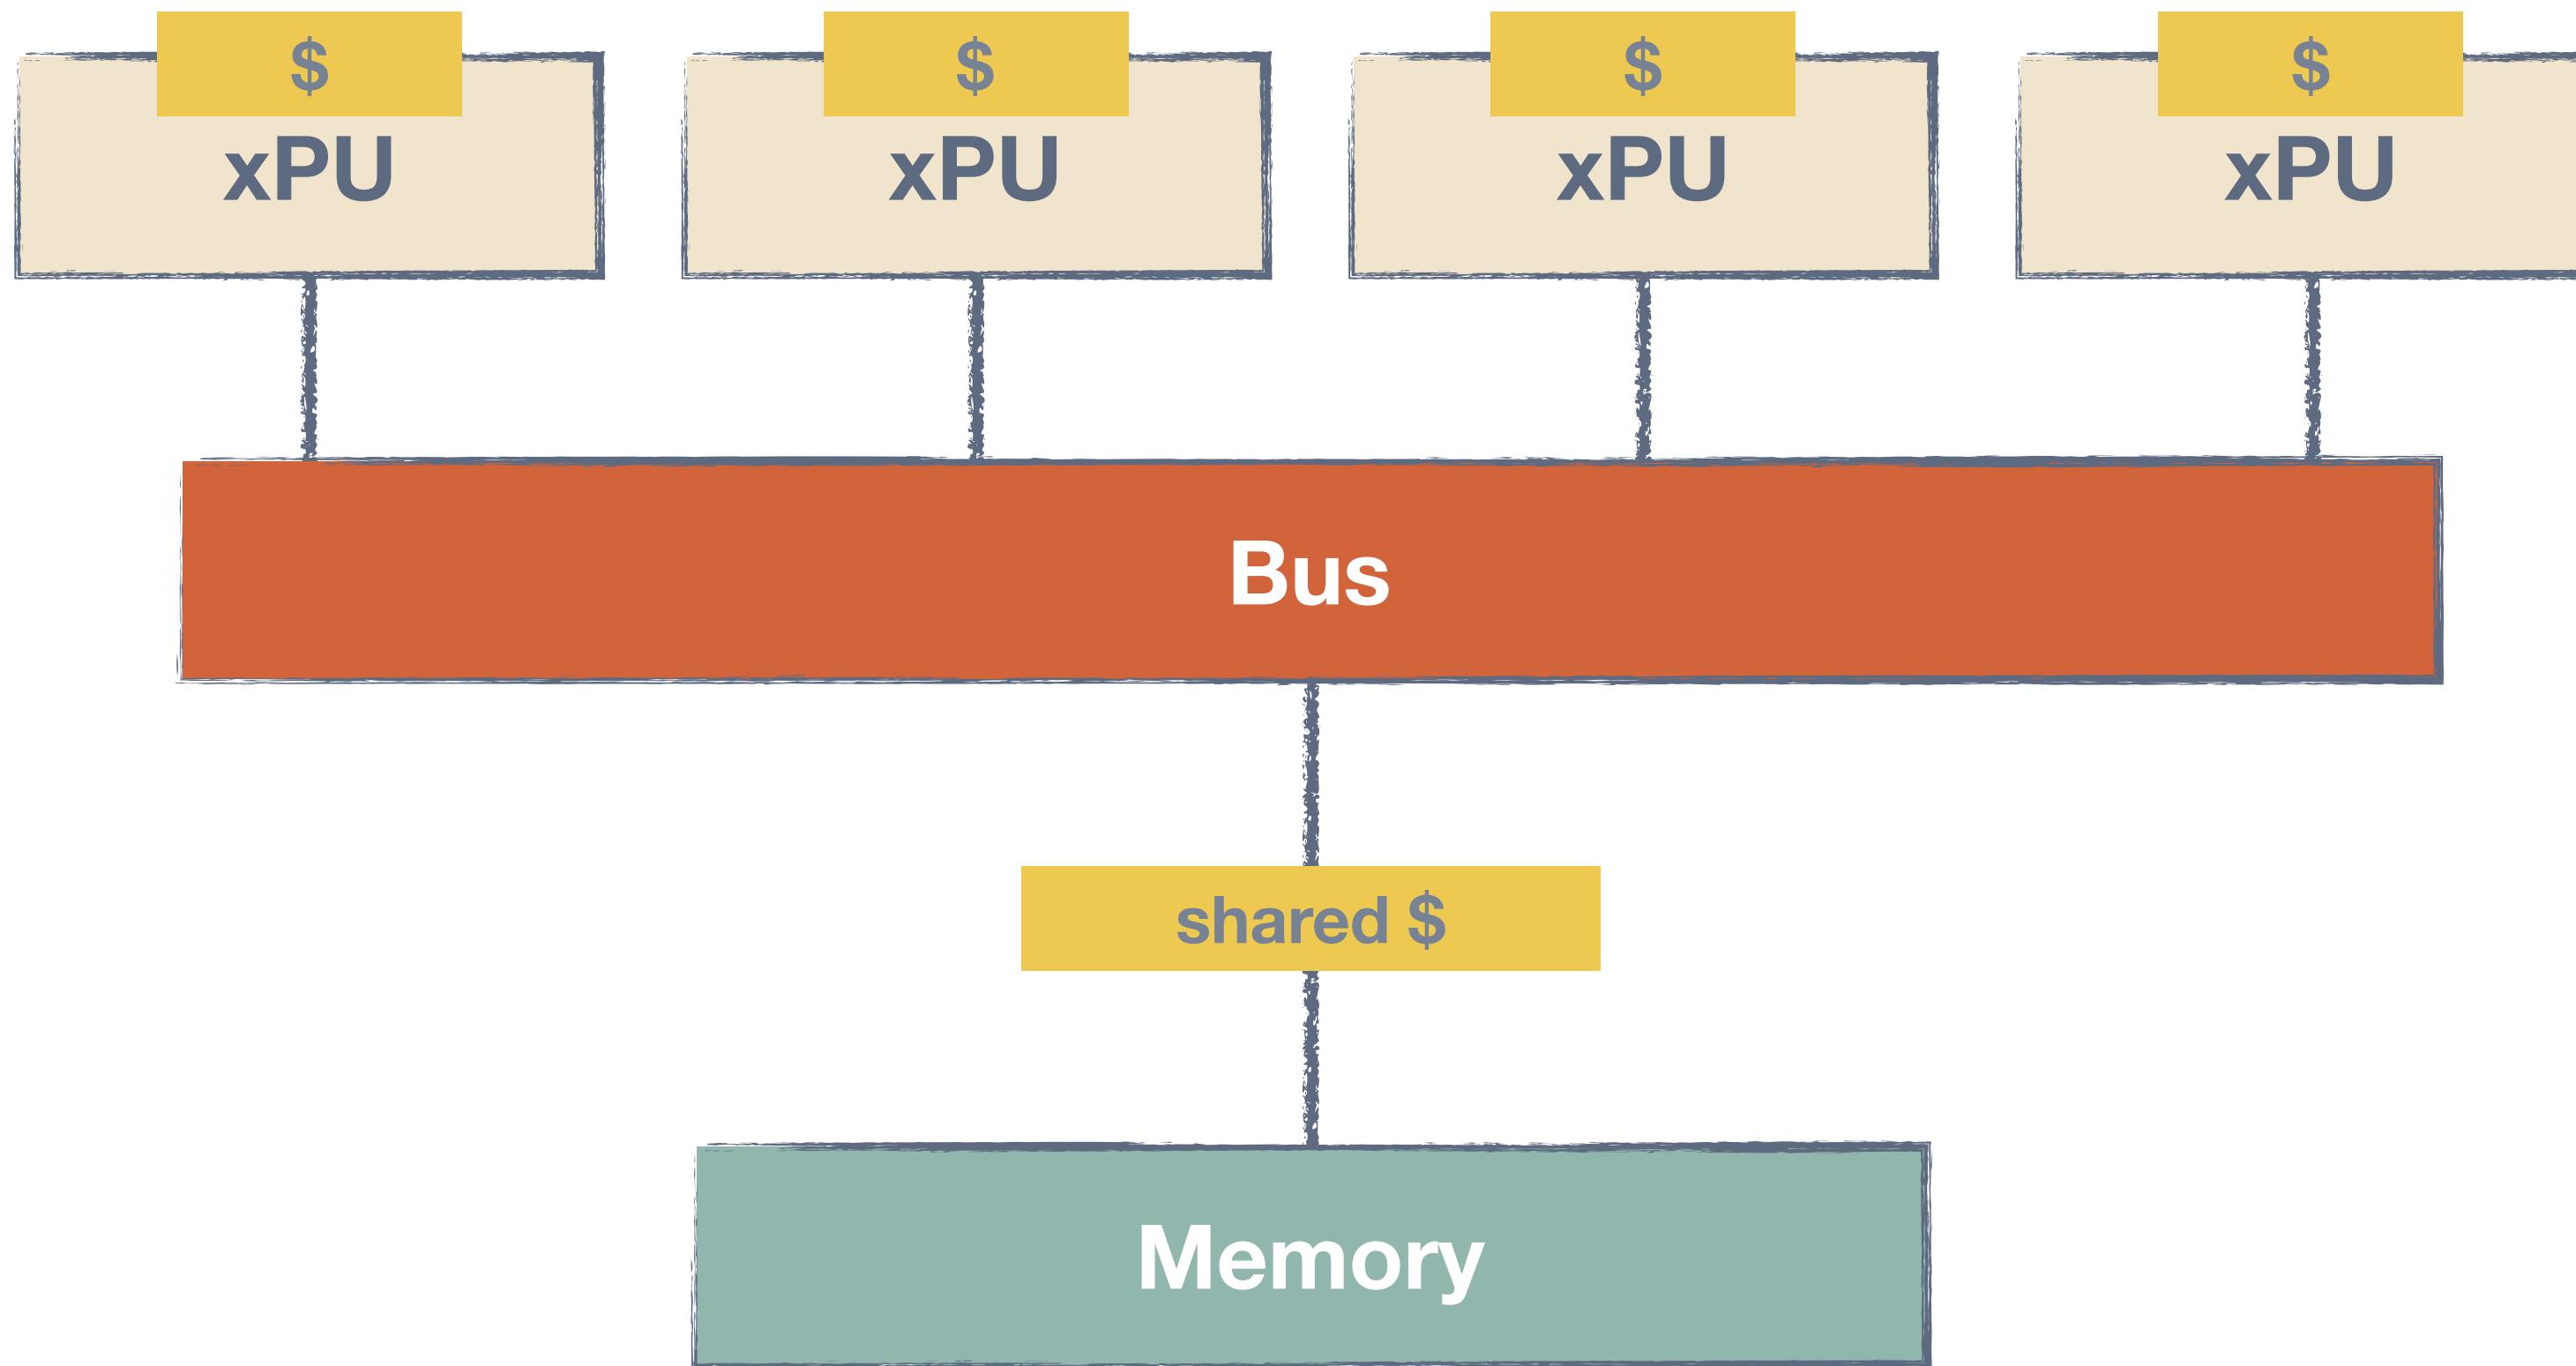
Speedup vs Efficiency

A fixed problem size limits the maximum speedup. But given more parallel hardware, one might choose instead to solve a larger problem. Scaling the problem is called **weak scaling** (vs. strong). A suitable measure in this case is not speedup but **parallel efficiency**, which we hope approaches a constant as $p \rightarrow \infty$.

$$E_p(n) = \frac{S_p(n)}{p}$$

$$E_p(n) = \frac{W(n)}{p \cdot T_p(n)}$$

$$E_p(n) \geq \frac{1}{\frac{D(n)}{W(n)}p + 1}$$

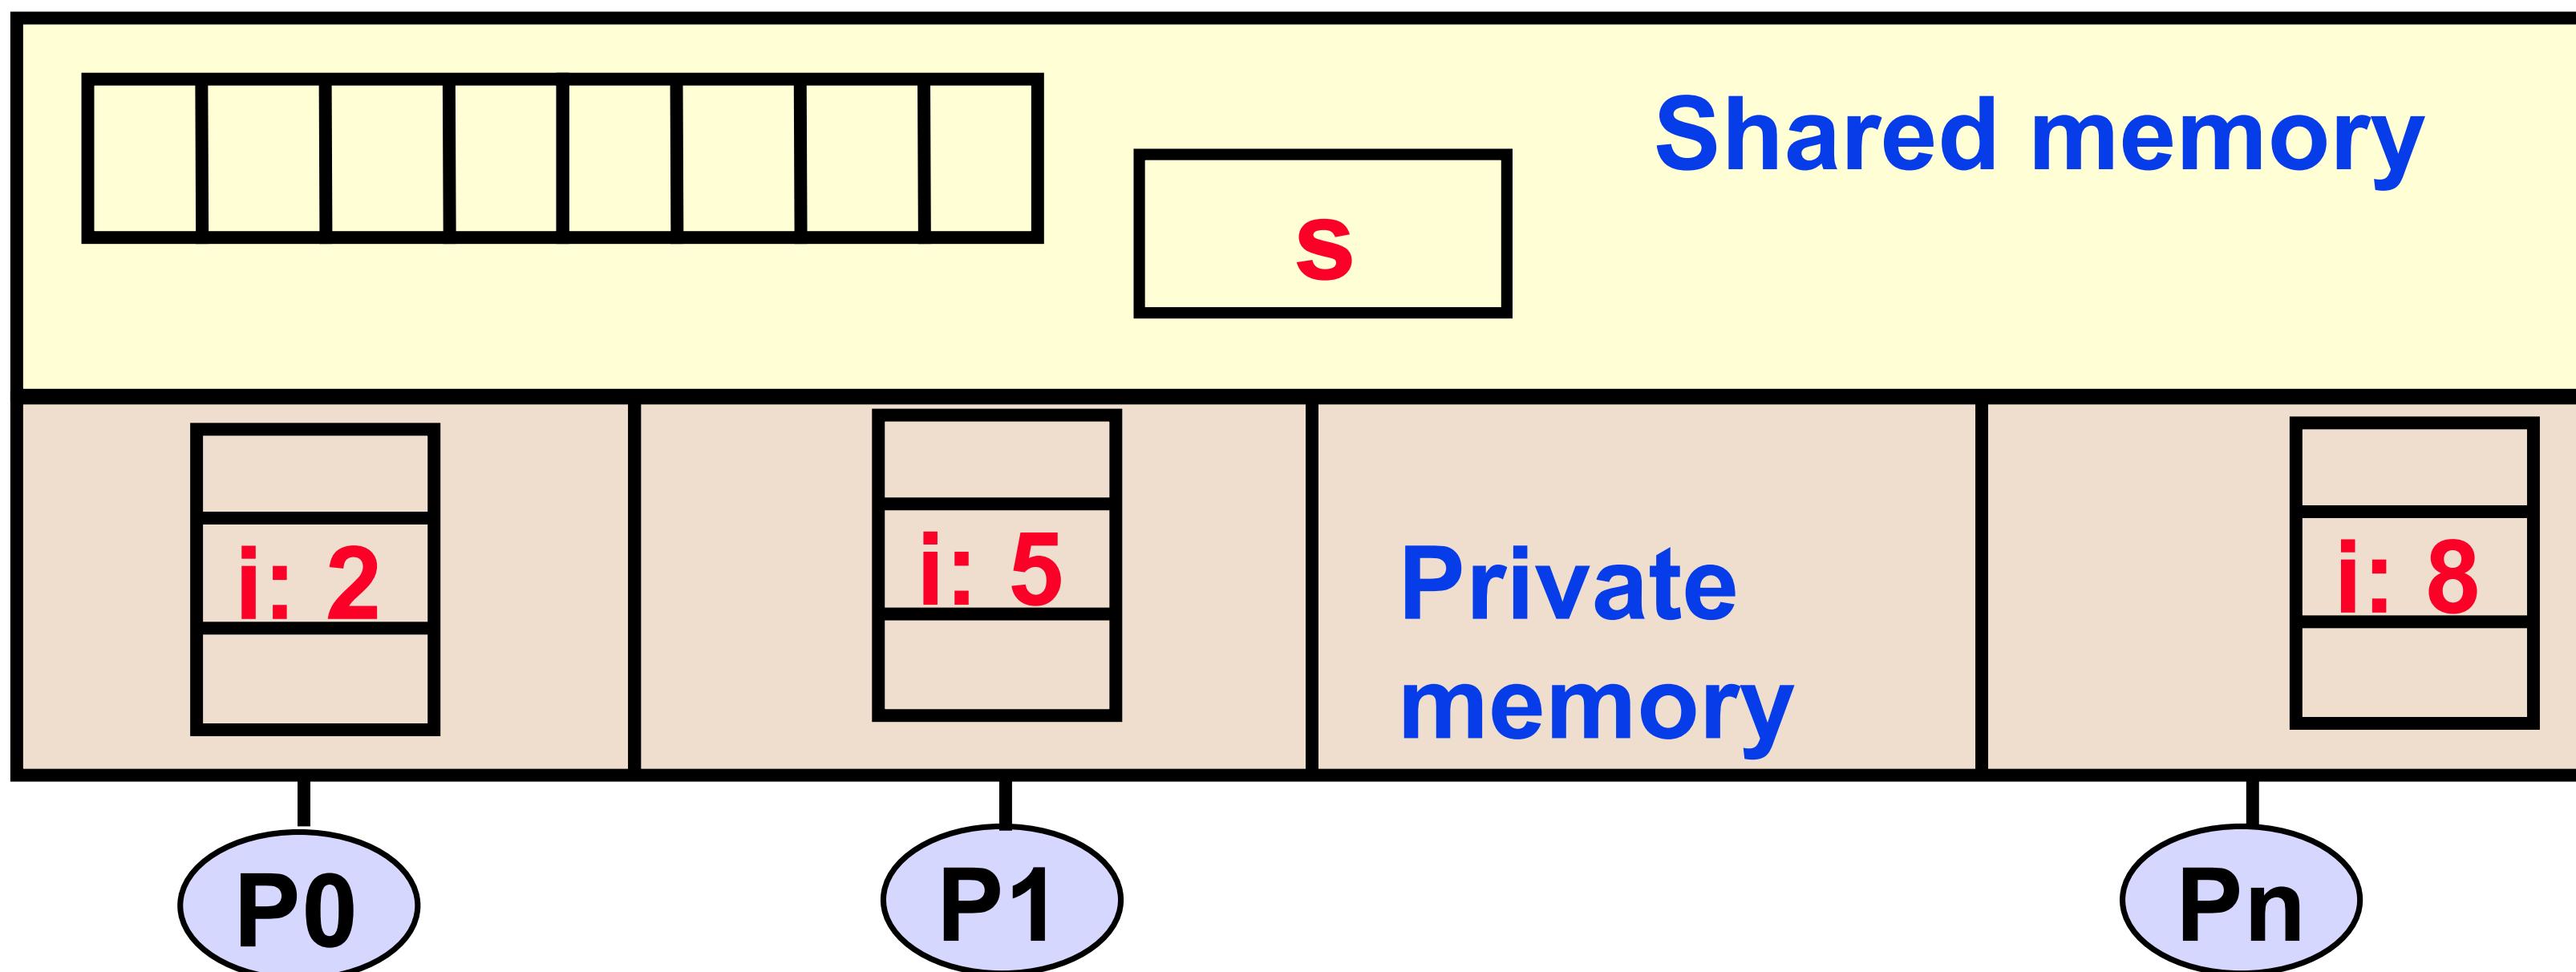


SHARED MEMORY

All processors are connected to a shared memory.
Advantage: Uniform memory access (UMA)
Disadvantage: Difficult to scale to large number of processors

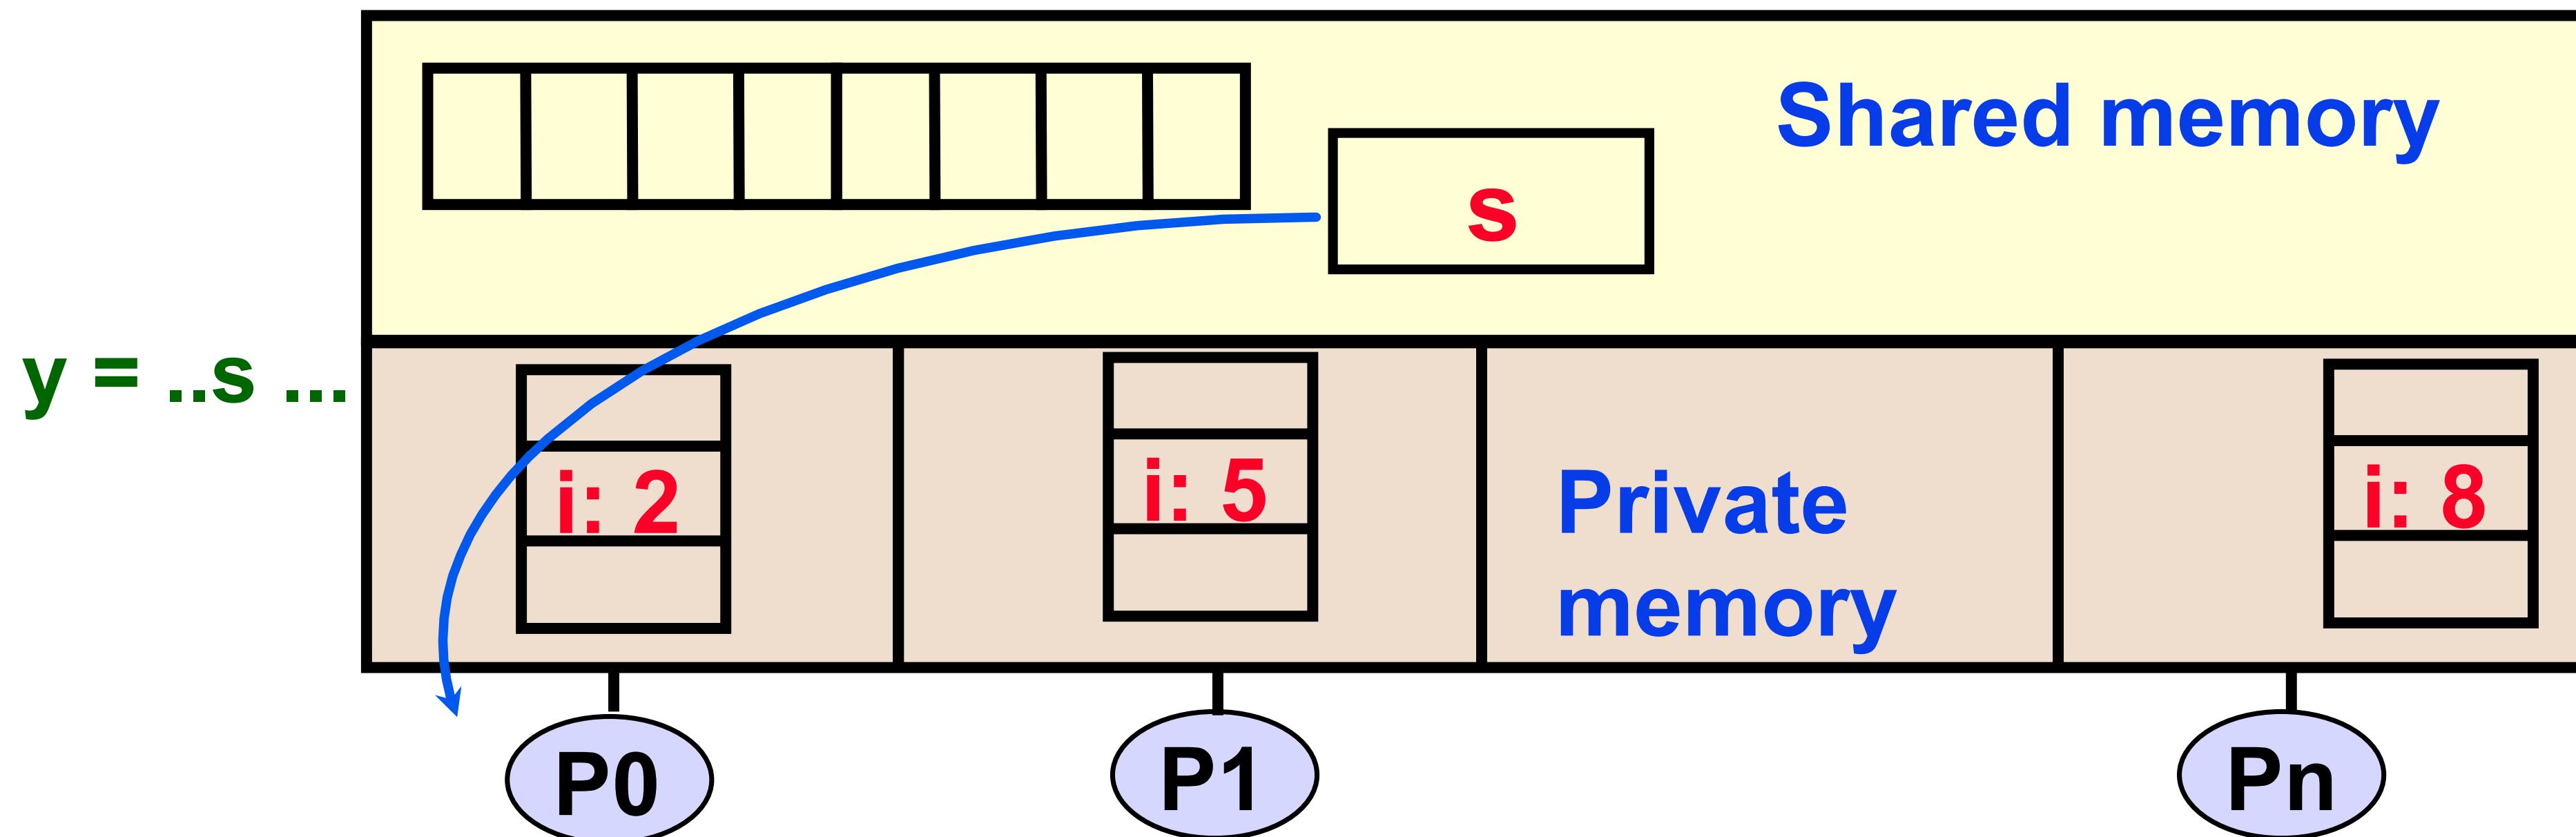
Shared Memory Model

- ▶ Program is a collection of threads of control
- ▶ Threads **communicate implicitly** by reading and writing shared variables
- ▶ Threads coordinate by synchronizing on shared variables



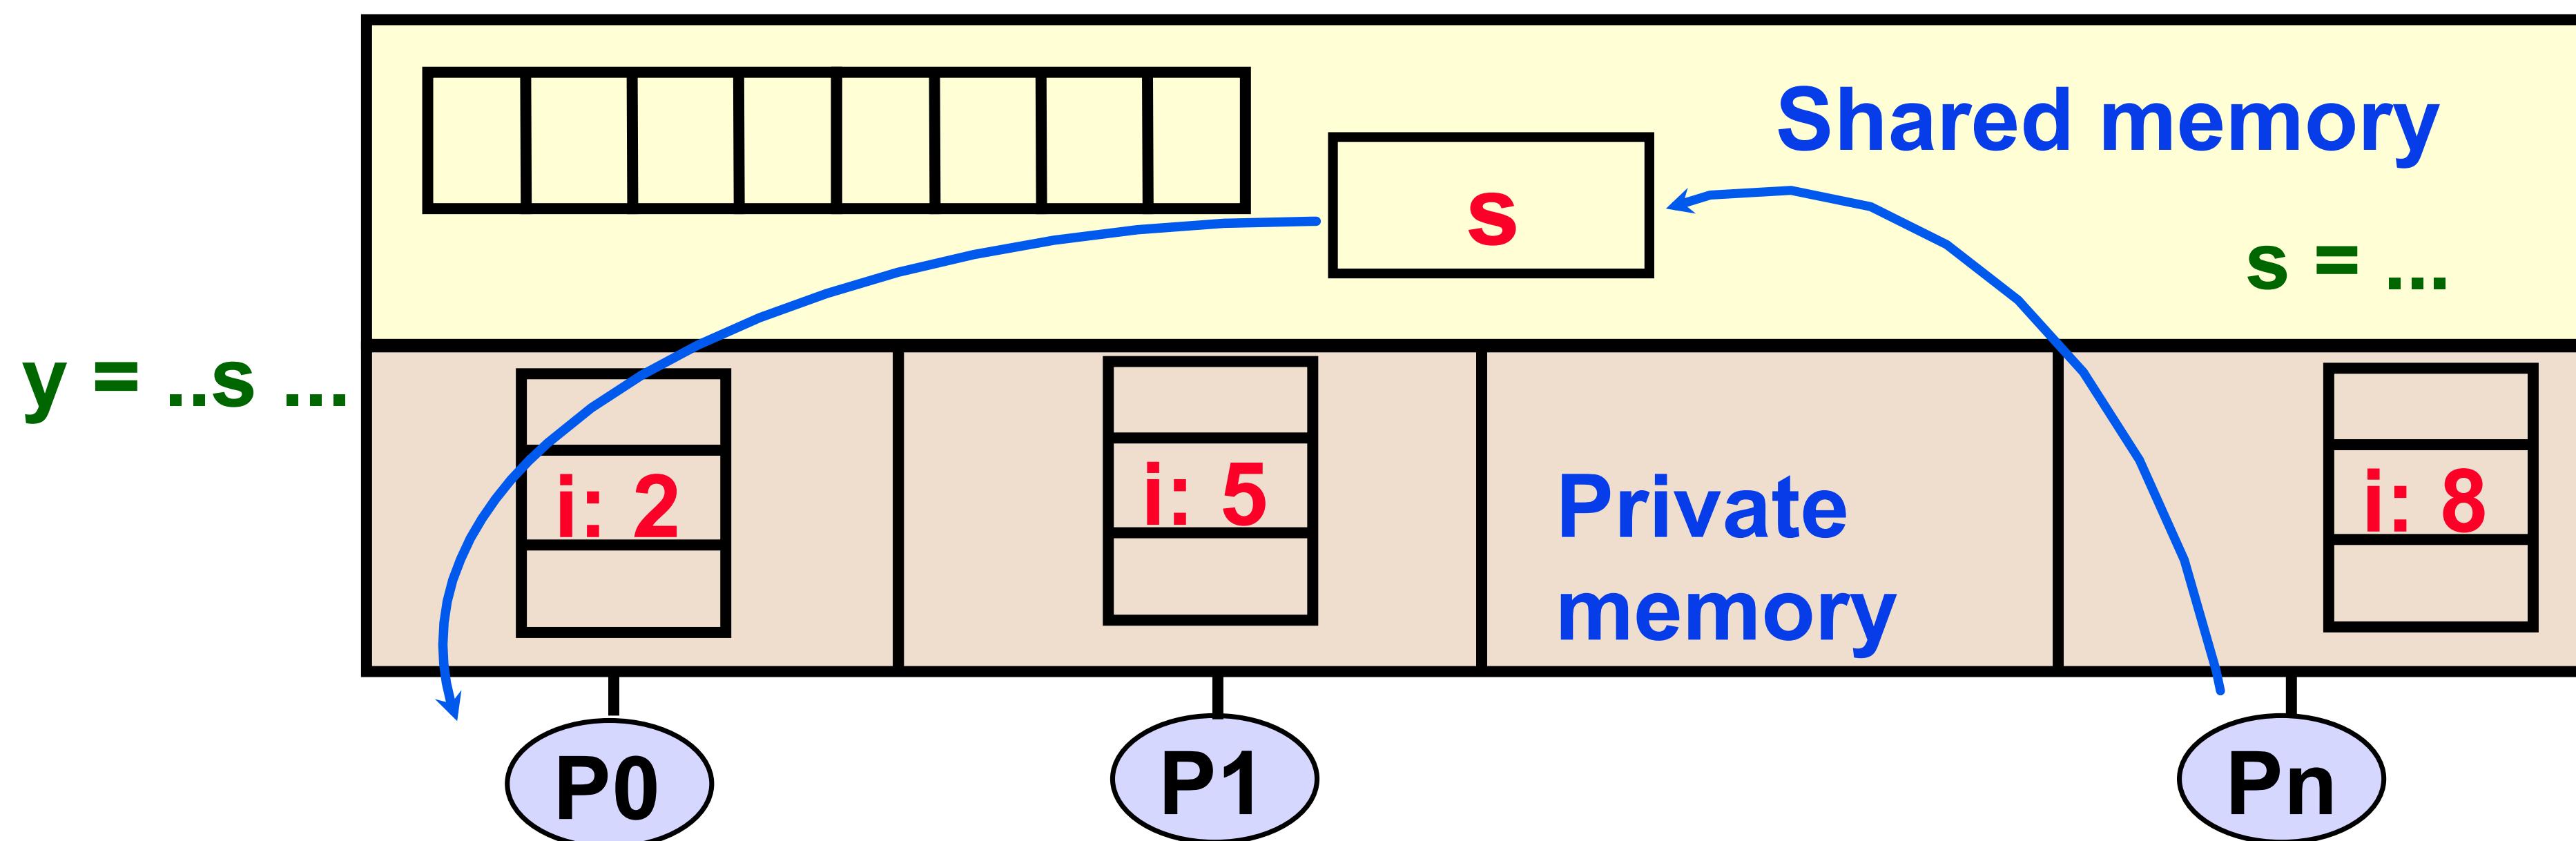
Shared Memory Model

- ▶ Program is a collection of threads of control
- ▶ Threads **communicate implicitly** by reading and writing shared variables
- ▶ Threads coordinate by synchronizing on shared variables



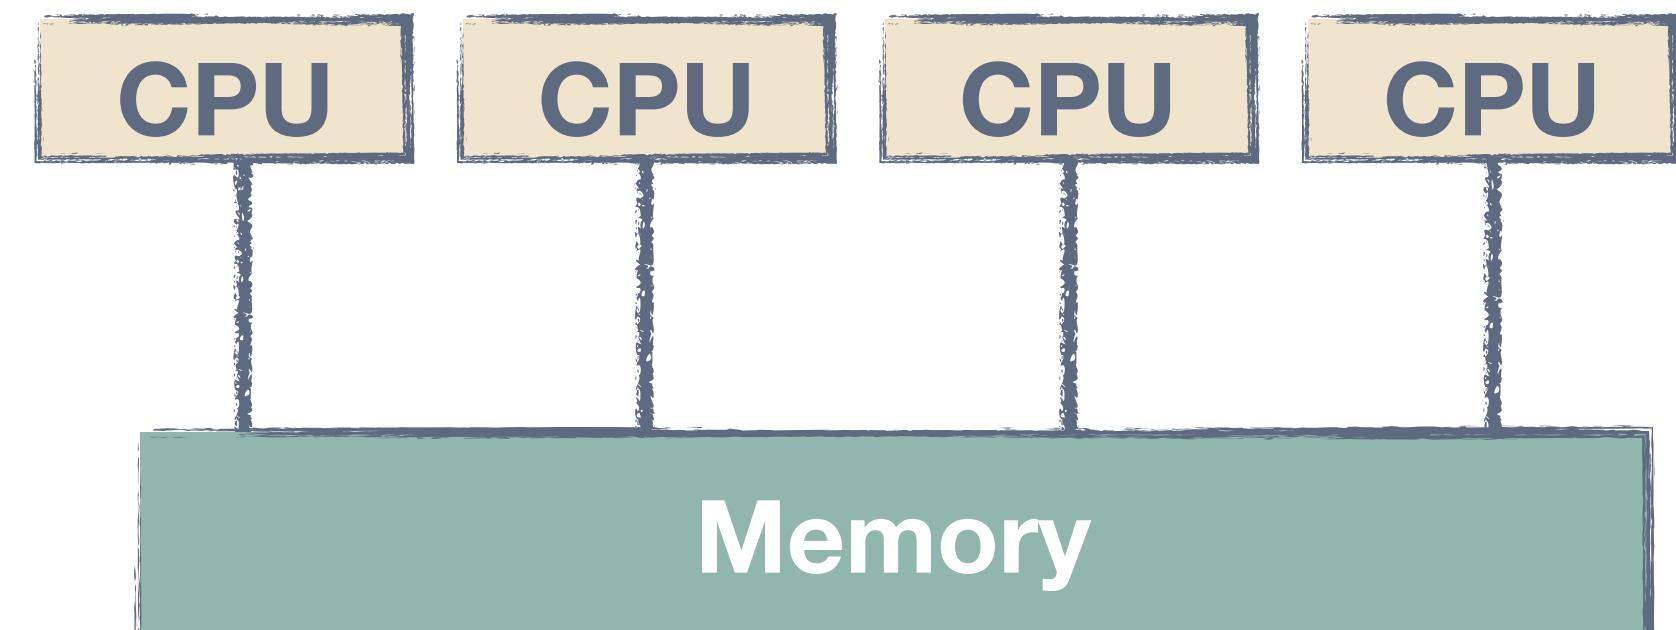
Shared Memory Model

- ▶ Program is a collection of threads of control
- ▶ Threads **communicate implicitly** by reading and writing shared variables
- ▶ Threads coordinate by synchronizing on shared variables

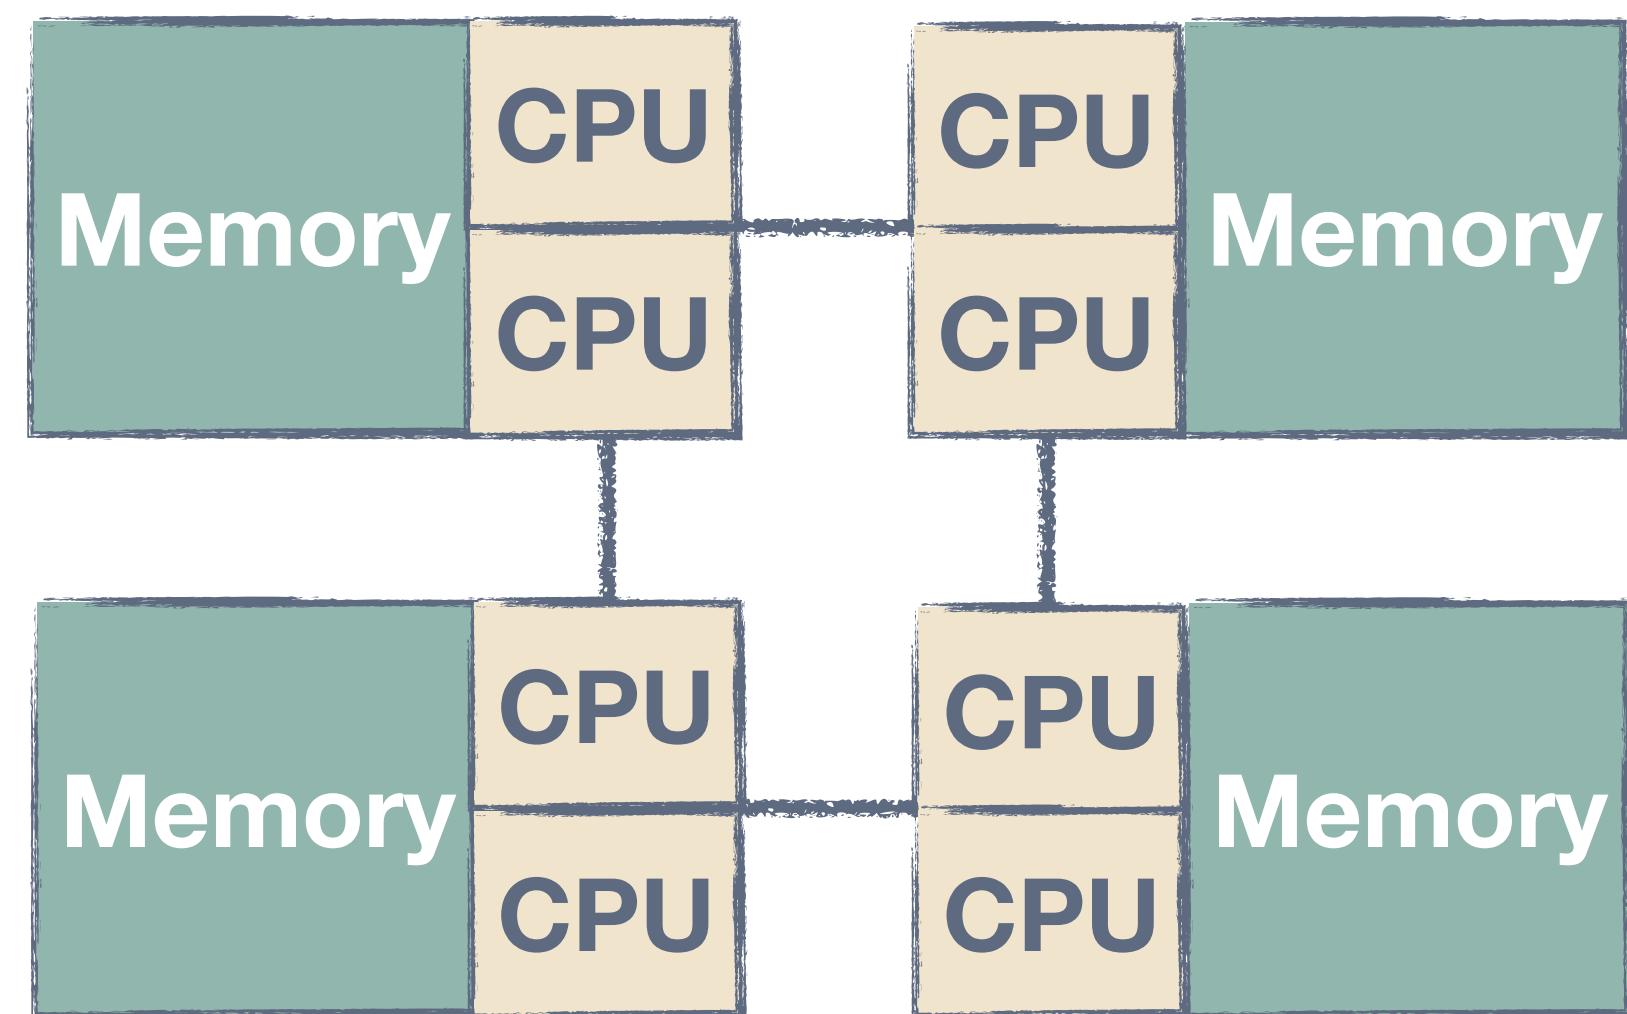


OpenMP programming model

- ▶ OpenMP is designed for shared memory machines. The underlying architecture can be shared memory **UMA** or **NUMA**.
- ▶ Comprised of three components:
 - ▶ Compiler directives
 - ▶ Library routines
 - ▶ Environment variables
- ▶ Parallelism through the use of threads
- ▶ Explicit programming model, offering the programmer full control over parallelization

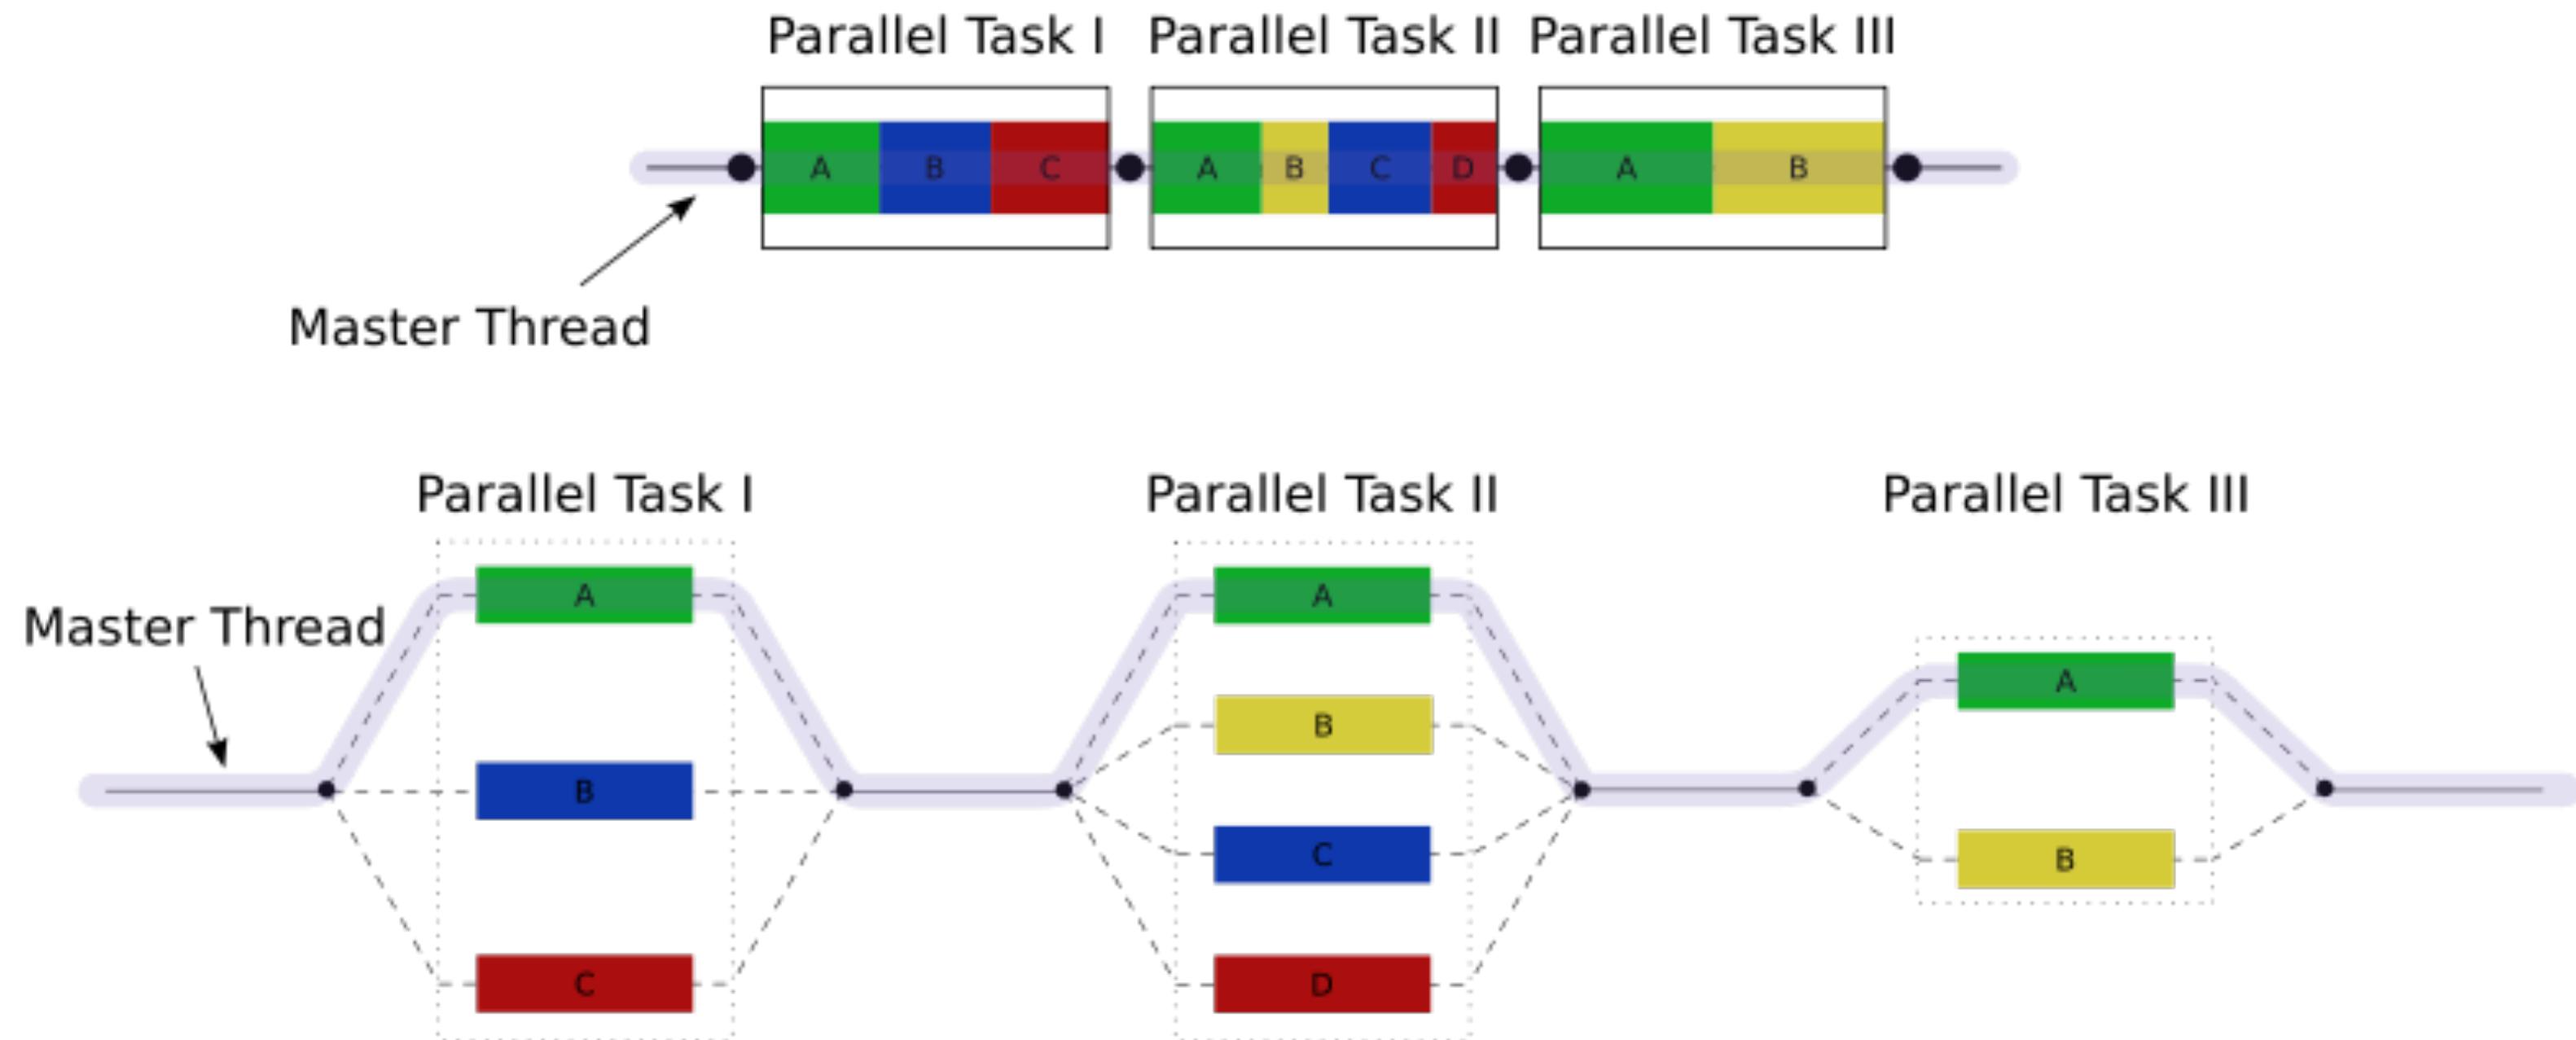


UNIFORM MEMORY Access (UMA)

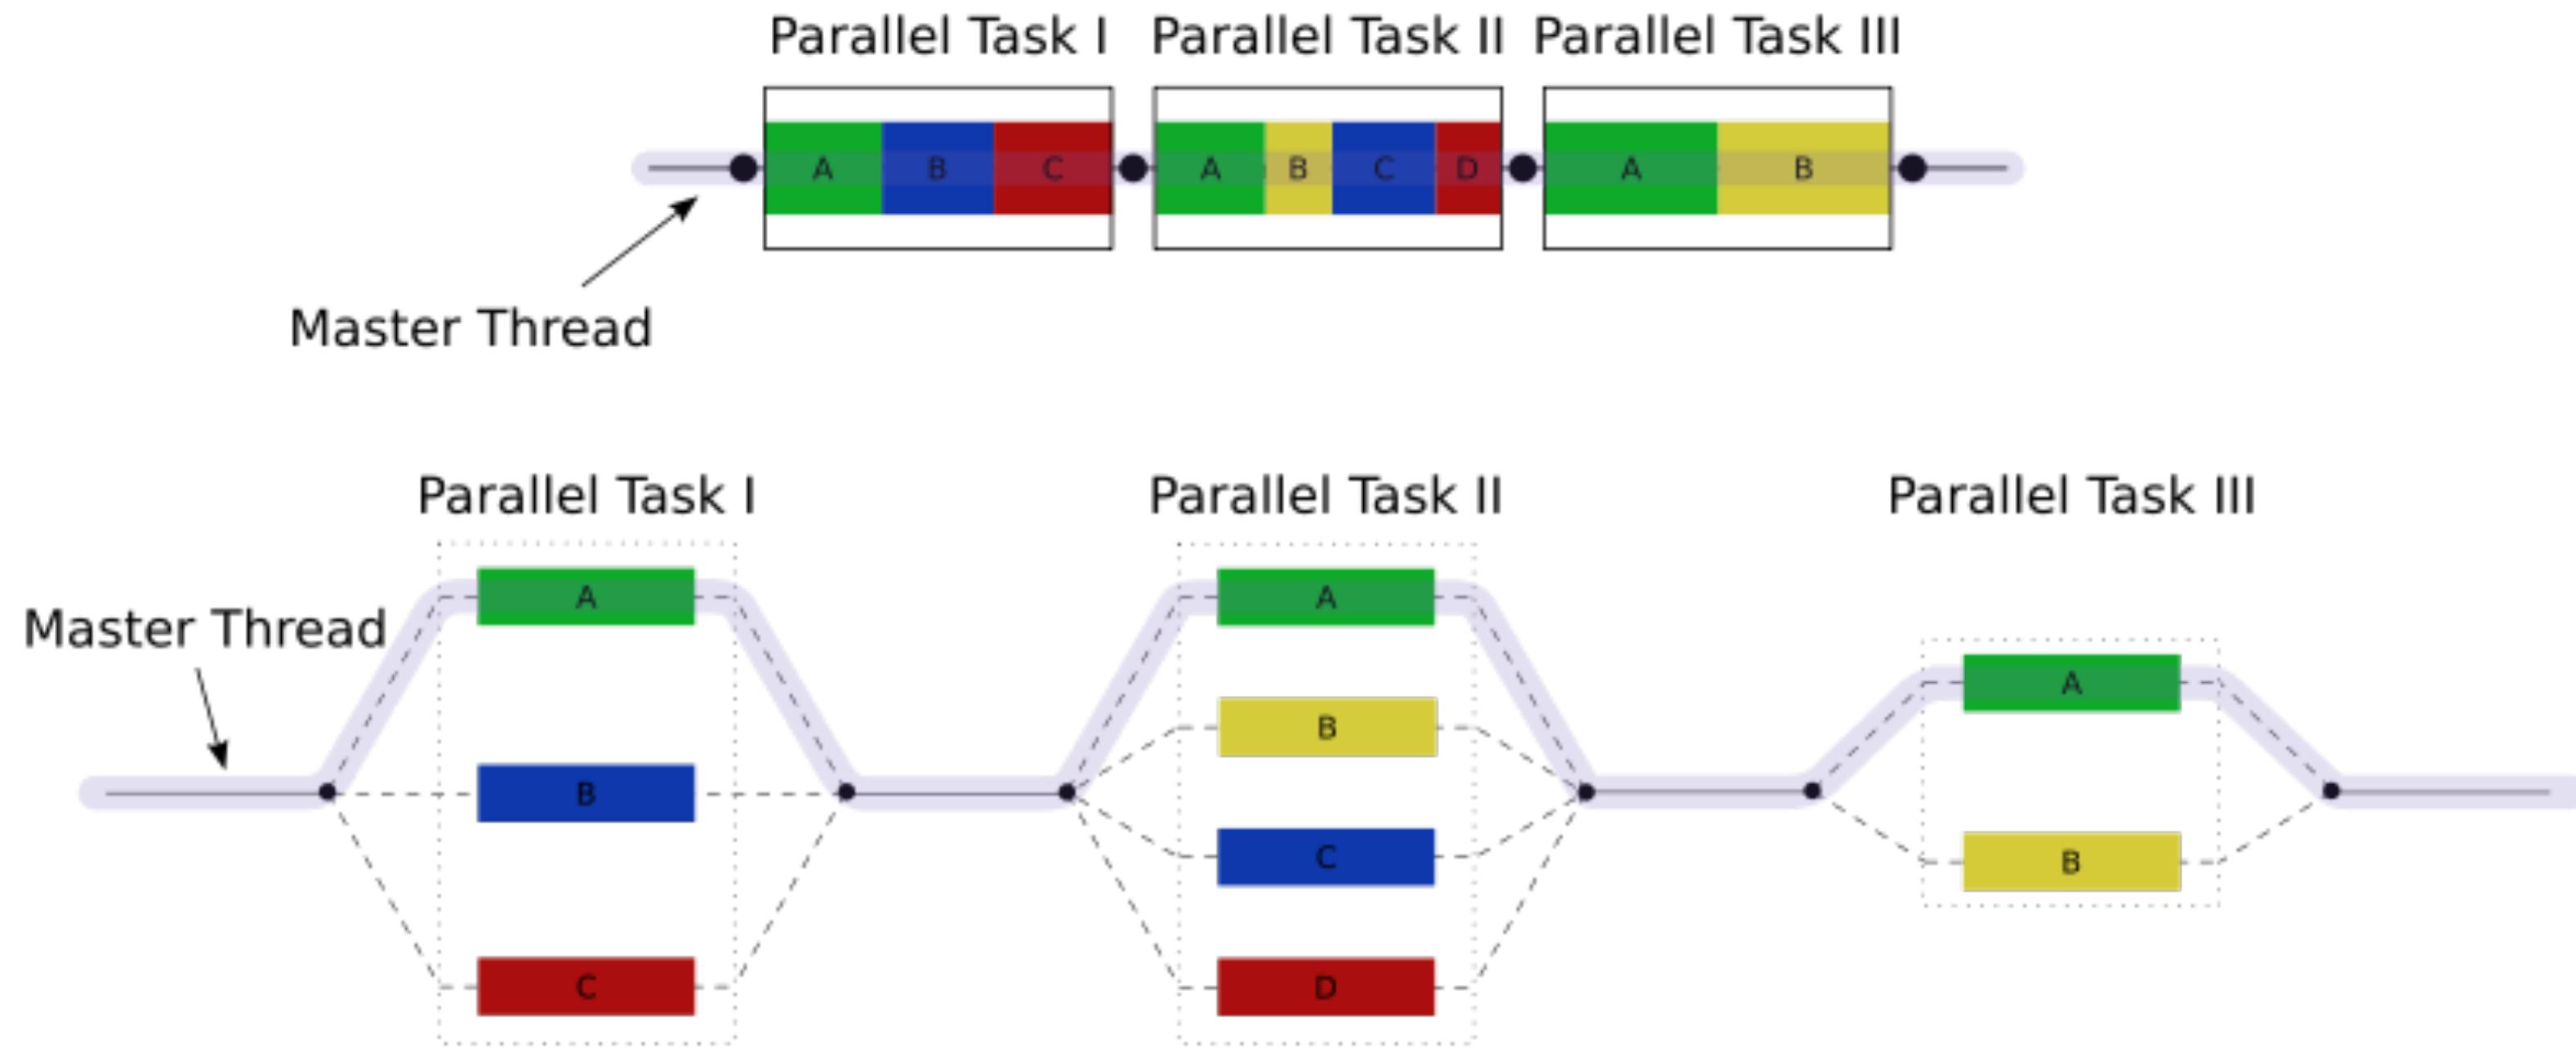


NON-UNIFORM MEMORY Access (NUMA)

Fork-Join model



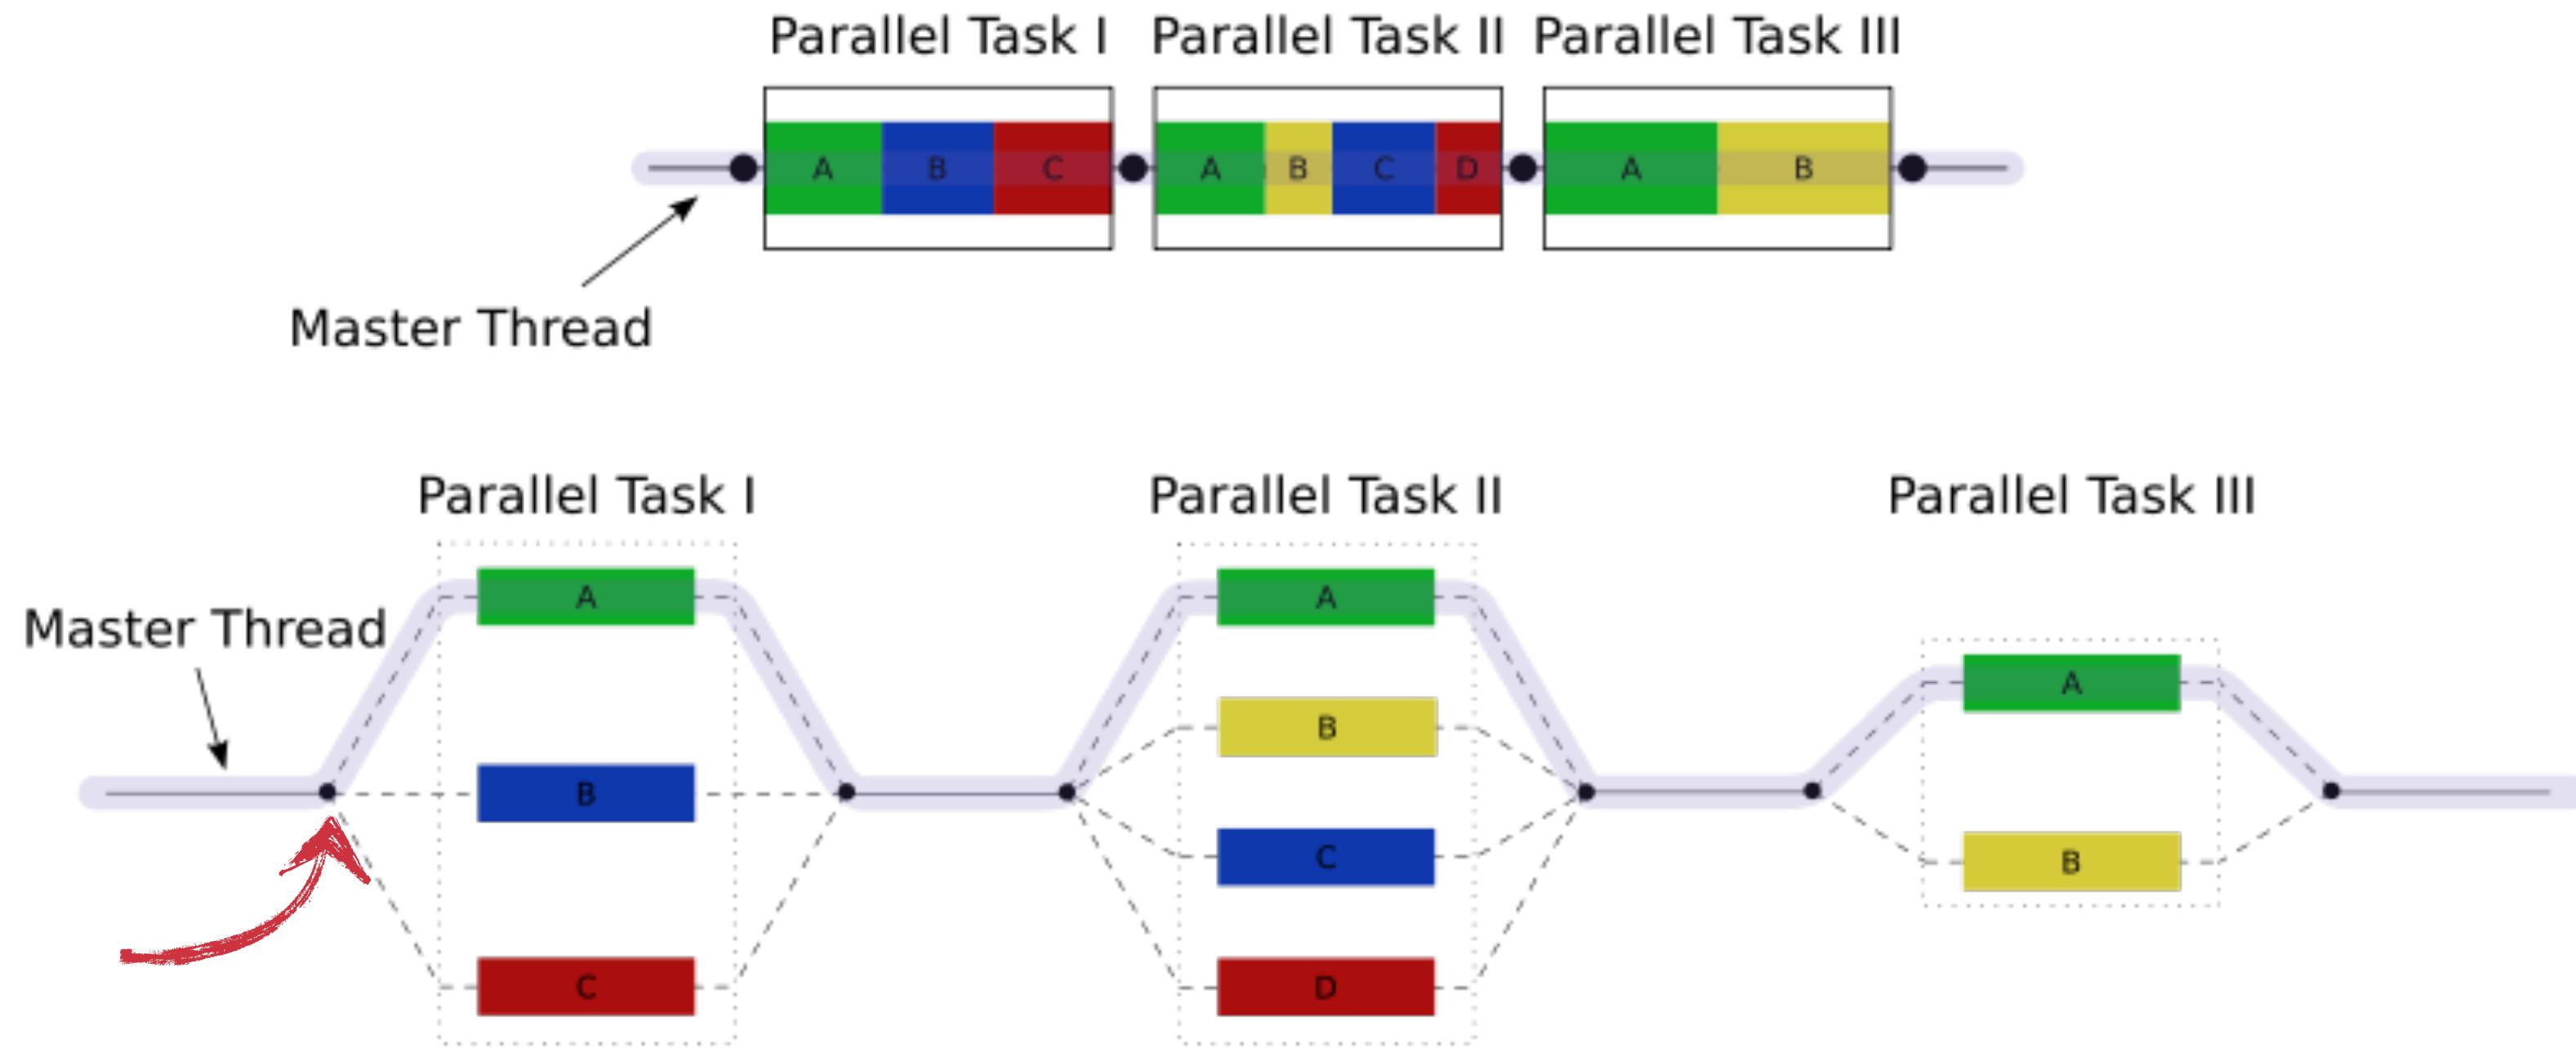
Fork-Join model



FORK: master thread then creates a **team** of parallel threads.

JOIN: team of threads execute the statements in the parallel region, then **synchronize** and terminate.

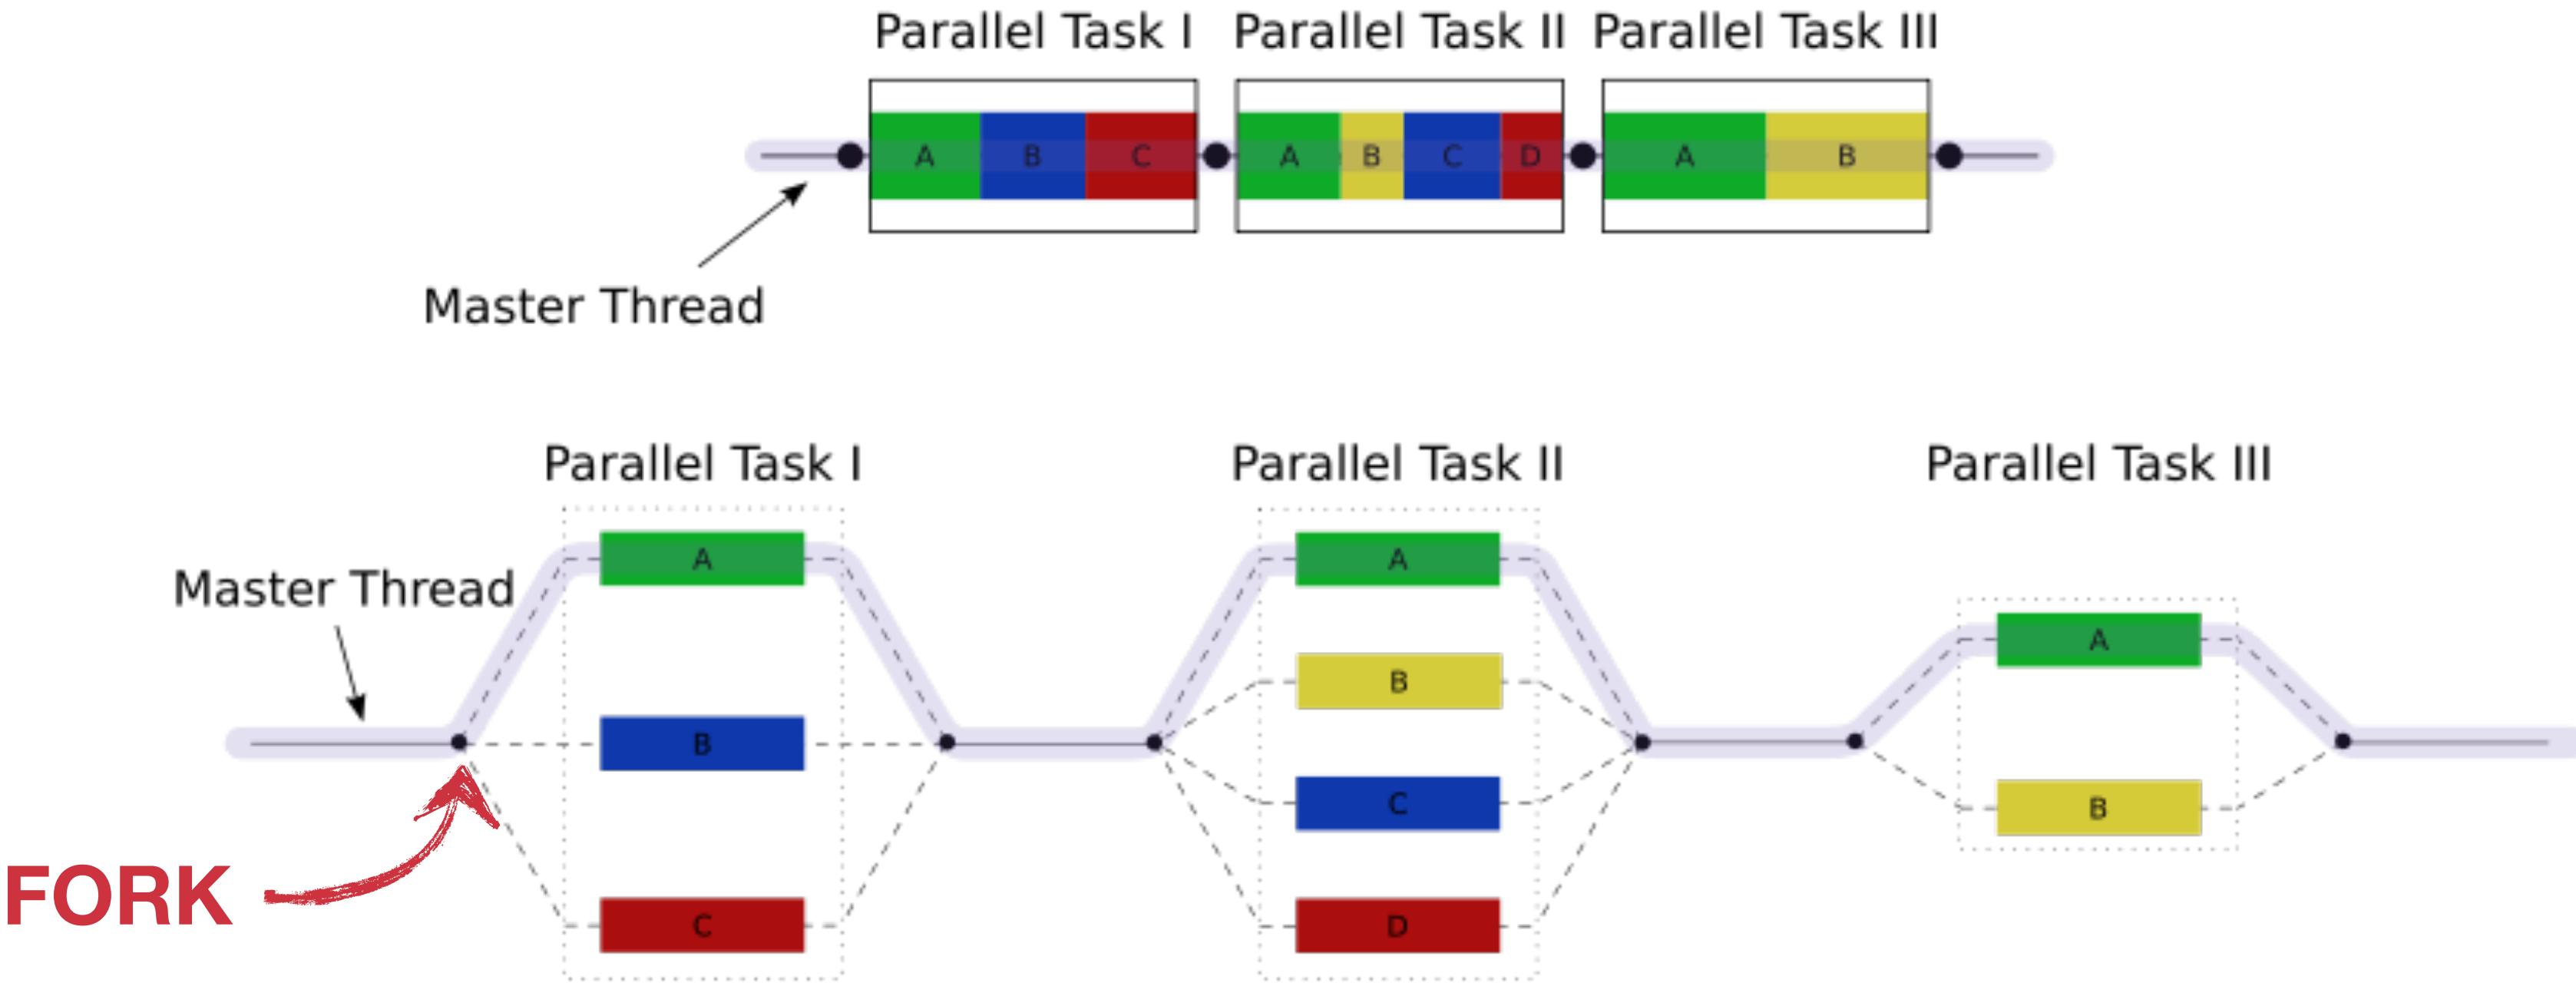
Fork-Join model



FORK: master thread then creates a **team** of parallel threads.

JOIN: team of threads execute the statements in the parallel region, then **synchronize** and terminate.

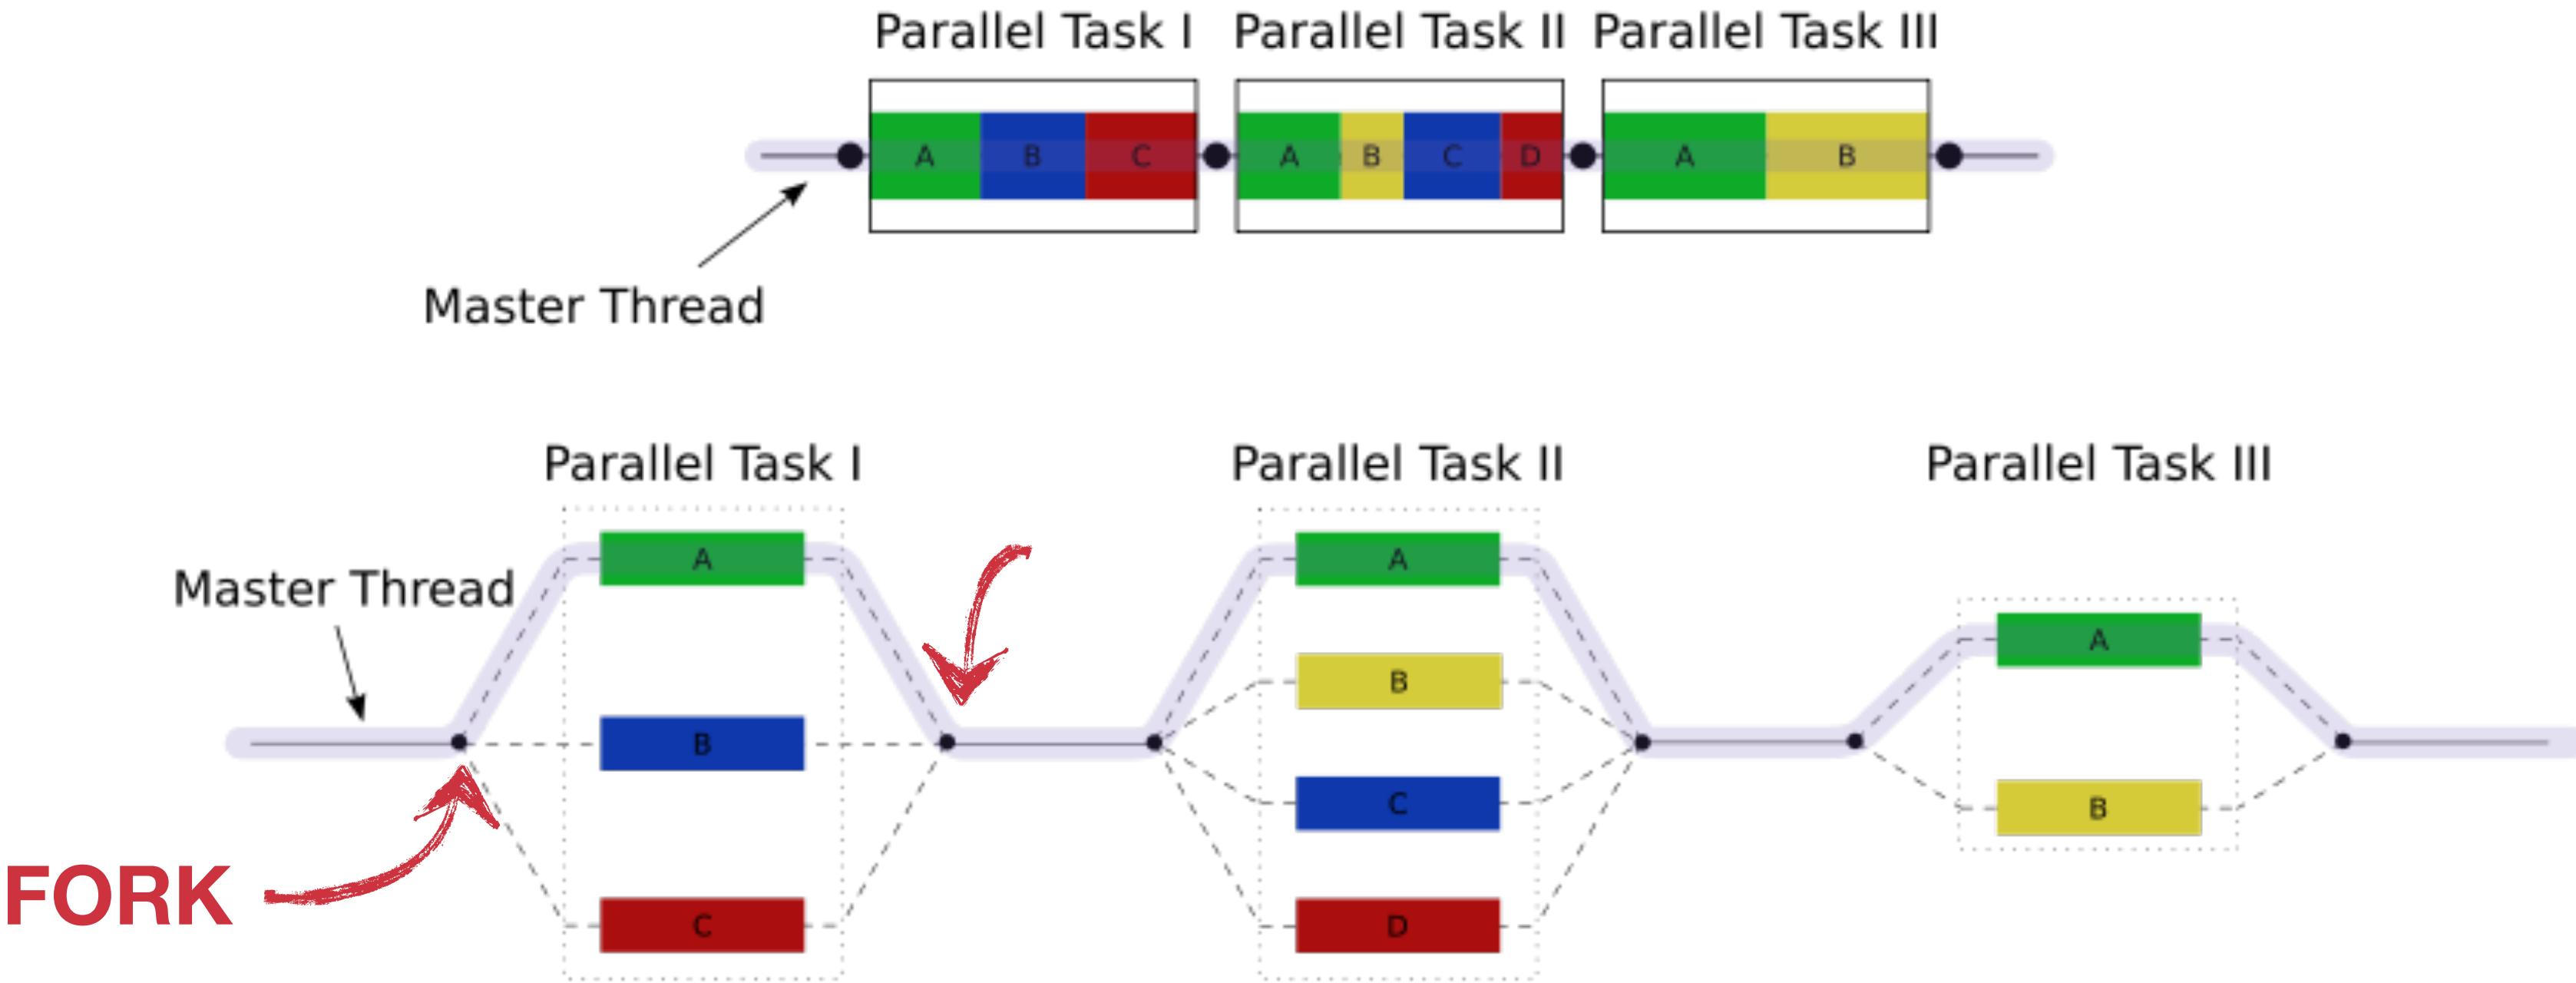
Fork-Join model



FORK: master thread then creates a **team** of parallel threads.

JOIN: team of threads execute the statements in the parallel region, then **synchronize** and terminate.

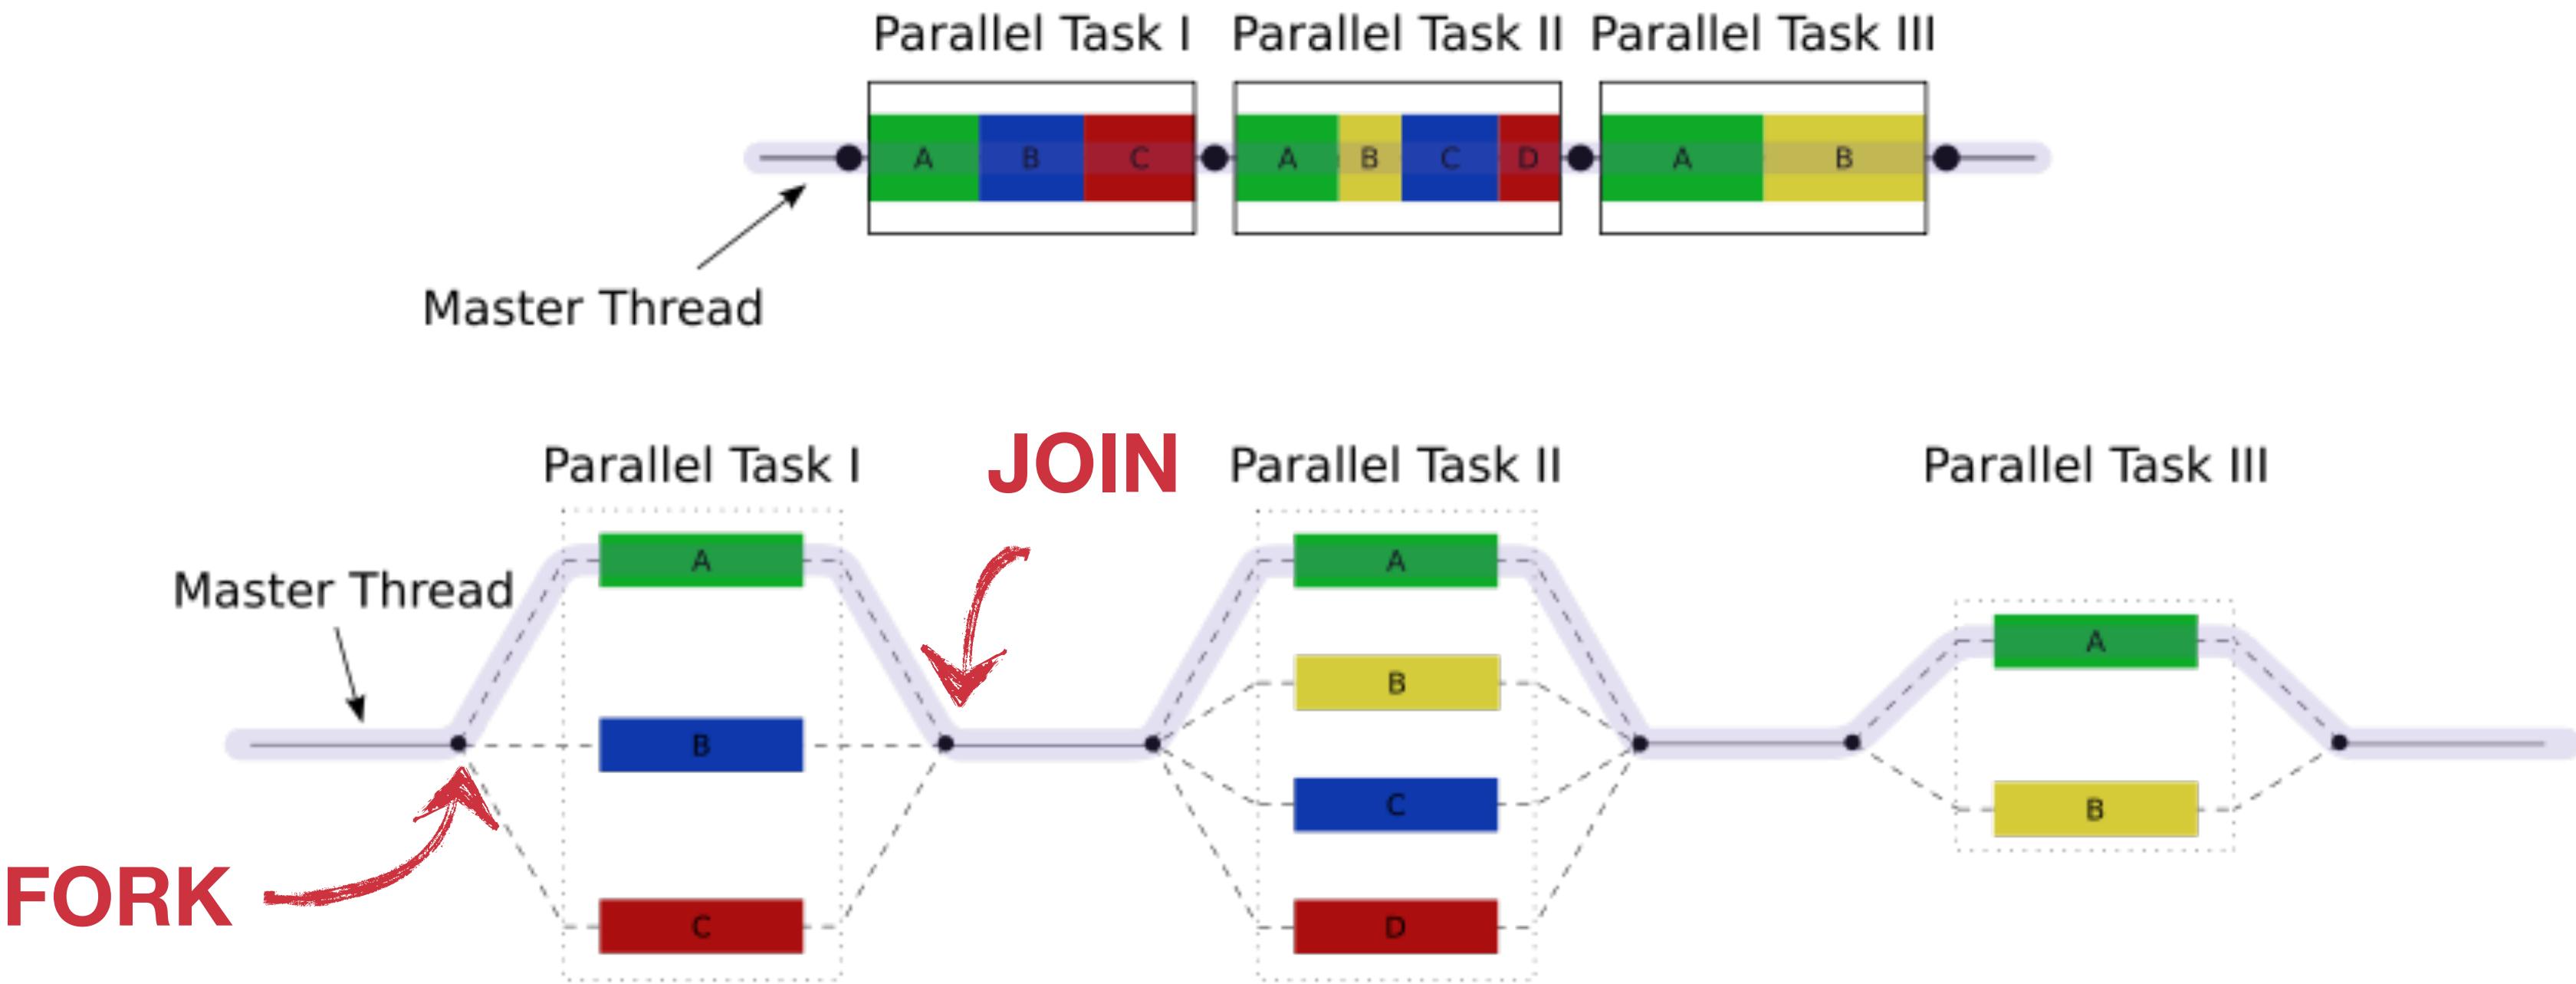
Fork-Join model



FORK: master thread then creates a **team** of parallel threads.

JOIN: team of threads execute the statements in the parallel region, then **synchronize** and terminate.

Fork-Join model



FORK: master thread then creates a **team** of parallel threads.

JOIN: team of threads execute the statements in the parallel region, then **synchronize** and terminate.

“Hello World” example

```
int main()
{
    printf ("hello, world!\n"); // Execute in parallel

    return 0;
}
```

“Hello World” example

```
#include <omp.h>

int main()
{
    printf ("hello, world!\n"); // Execute in parallel

    return 0;
}
```

“Hello World” example

```
#include <omp.h>

int main()
{
    #pragma omp parallel
    {
        printf ("hello, world!\n"); // Execute in parallel
    } // Implicit barrier/join
    return 0;
}
```

“Hello World” example

```
#include <omp.h>

int main()
{
    omp_set_num_threads(16); // Can also use OMP_NUM_THREADS
                            // environment variable

#pragma omp parallel
{
    printf("hello, world!\n"); // Execute in parallel
} // Implicit barrier/join
return 0;
}
```

“Hello World” example

```
#include <omp.h>

int main()
{
    omp_set_num_threads(16); // Can also use OMP_NUM_THREADS
                            // environment variable

#pragma omp parallel
{
    printf("hello, world!\n"); // Execute in parallel
} // Implicit barrier/join
return 0;
}
```

Compiling:

gcc -fopenmp
icc -openmp
pgcc -mp

Parallel loops

```
for (i = 0; i < n; ++i) {  
    a[i] += i;  
}
```

Parallel loops

```
for (i = 0; i < n; ++i) {  
    a[i] += i;  
}
```

```
#pragma omp parallel  
{  
#pragma omp for shared (a,n) private (i)  
for (i = 0; i < n; ++i) {  
    a[i] += i;  
} // Implicit barrier/join  
} // Implicit barrier/join
```

Parallel loops

```
for (i = 0; i < n; ++i) {  
    a[i] += i;  
}
```

Parallel loops

```
for (i = 0; i < n; ++i) {  
    a[i] += i;  
}
```

```
#pragma omp parallel for shared (a,n) private (i)  
for (i = 0; i < n; ++i) {  
    a[i] += i;  
} // Implicit barrier/join
```

Parallel loops with dependencies

```
sum = 0;  
for (i = 0; i < n; ++i) {  
    sum += i;  
}
```

Parallel loops with dependencies

```
sum = 0;  
for (i = 0; i < n; ++i) {  
    sum += i;  
}
```

```
#pragma omp parallel for shared (sum,n) private (i)  
for (i = 0; i < n; ++i) {  
    // Data race condition!  
    sum += i;  
}
```

Parallel loops with dependencies

```
sum = 0;  
for (i = 0; i < n; ++i) {  
    sum += i;  
}
```

```
#pragma omp parallel for shared (sum,n) private (i)  
for (i = 0; i < n; ++i) {  
    #pragma omp critical  
    sum += i;  
}
```

Parallel loops with dependencies

```
sum = 0;  
for (i = 0; i < n; ++i) {  
    sum += i;  
}
```

```
#pragma omp parallel for reduction (+:sum)  
for (i = 0; i < n; ++i) {  
    sum += i;  
}
```

Parallel loop scheduling

- ▶ Static: *chunk* iterations per thread and statically assigned
#pragma omp parallel for schedule (static, chunk)
- ▶ Dynamic: *chunk* iterations per thread and dynamically assigned
#pragma omp parallel for schedule (dynamic, chunk)
- ▶ Guided: *chunk* iterations per thread initially, decreases with each allocation and dynamically assigned
#pragma omp parallel for schedule (guided, chunk)
- ▶ Runtime: set environment variable **OMP_SCHEDULE**

Synchronization primitive 1: Single

```
#pragma omp parallel shared (a, n) private (i)
{
    #pragma omp single
    for (i = 0; i < n; ++i) {
        a[i] += i;
    } // Implicit barrier/join
} // Implicit barrier/join
```

Synchronization primitive 1: Single

```
#pragma omp parallel shared (a, n) private (i)
{
    #pragma omp single nowait
    for (i = 0; i < n; ++i) {
        a[i] += i;
    }
} // Implicit barrier/join
```

Synchronization primitive 2: Master

```
#pragma omp parallel shared (a, n) private (i)
{
    #pragma omp master
    for (i = 0; i < n; ++i) {
        a[i] += i;
    } // No Implicit barrier/join
} // Implicit barrier/join
```

Synchronization primitive 3: Barrier

```
#pragma omp parallel shared (a, b, n) private (i)
{
    #pragma omp for nowait
    for (i = 0; i < n; ++i) {
        a[i] += i;
    }

    #pragma omp for
    for (i = 0; i < n; ++i) {
        b[i] = a[i] * 2;
    }
} // Implicit barrier/join
```

Synchronization primitive 3: Barrier

```
#pragma omp parallel shared (a, b, n) private (i)
{
    #pragma omp for nowait
    for (i = 0; i < n; ++i) {
        a[i] += i;
    }
    #pragma omp barrier
    #pragma omp for
    for (i = 0; i < n; ++i) {
        b[i] = a[i] * 2;
    }
} // Implicit barrier/join
```

Synchronization primitives

Master	#pragma omp master { }
Single	#pragma omp single { }
Critical	#pragma omp critical { }
Barrier	#pragma omp barrier
Locks	omp_set_lock(l) omp_unset_lock(l)

Tasking (OpenMP 3.0+)

```
// At the call site:
```

```
answer = fib (n);
```

```
// ...
```

```
int fib (int n) {  
    if (n <= G) fib_seq (n);  
    int f1, f2;
```

```
f1 = fib (n-1);  
f2 = fib (n-2);
```

```
return f1 + f2;
```

```
}
```

Tasking (OpenMP 3.0+)

```
// At the call site:  
#pragma omp parallel  
#pragma omp single nowait  
answer = fib (n);  
  
// ...  
  
int fib (int n) {  
    if (n <= G) fib_seq (n);  
    int f1, f2;  
  
    f1 = fib (n-1);  
    f2 = fib (n-2);  
  
    return f1 + f2;  
}
```

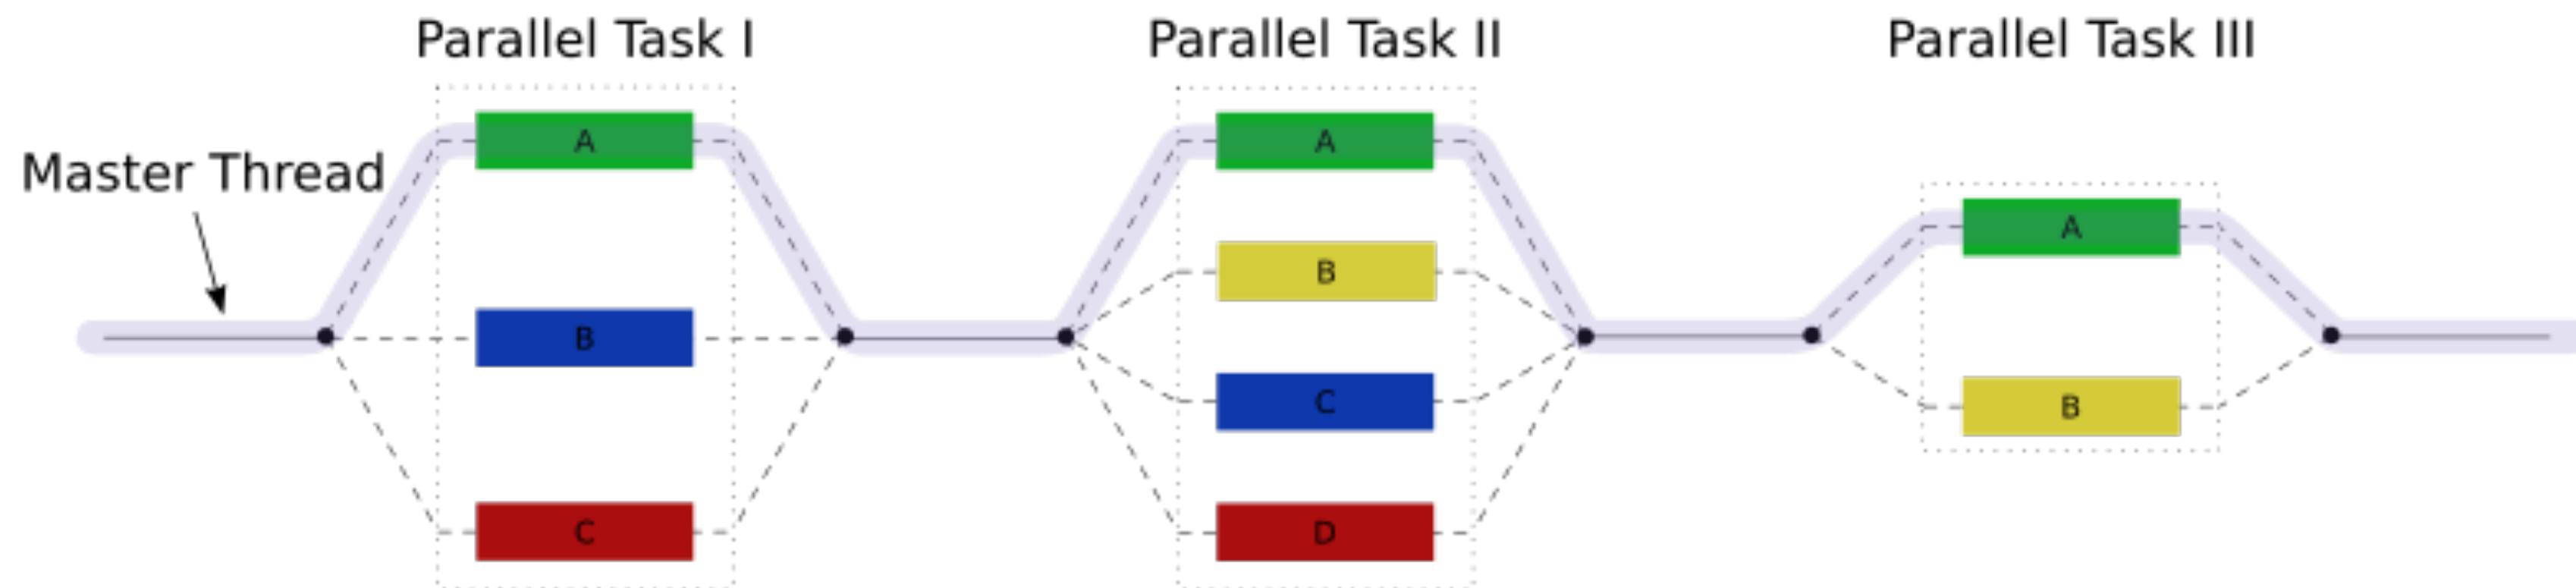
Tasking (OpenMP 3.0+)

```
// At the call site:  
#pragma omp parallel  
#pragma omp single nowait  
answer = fib (n);
```

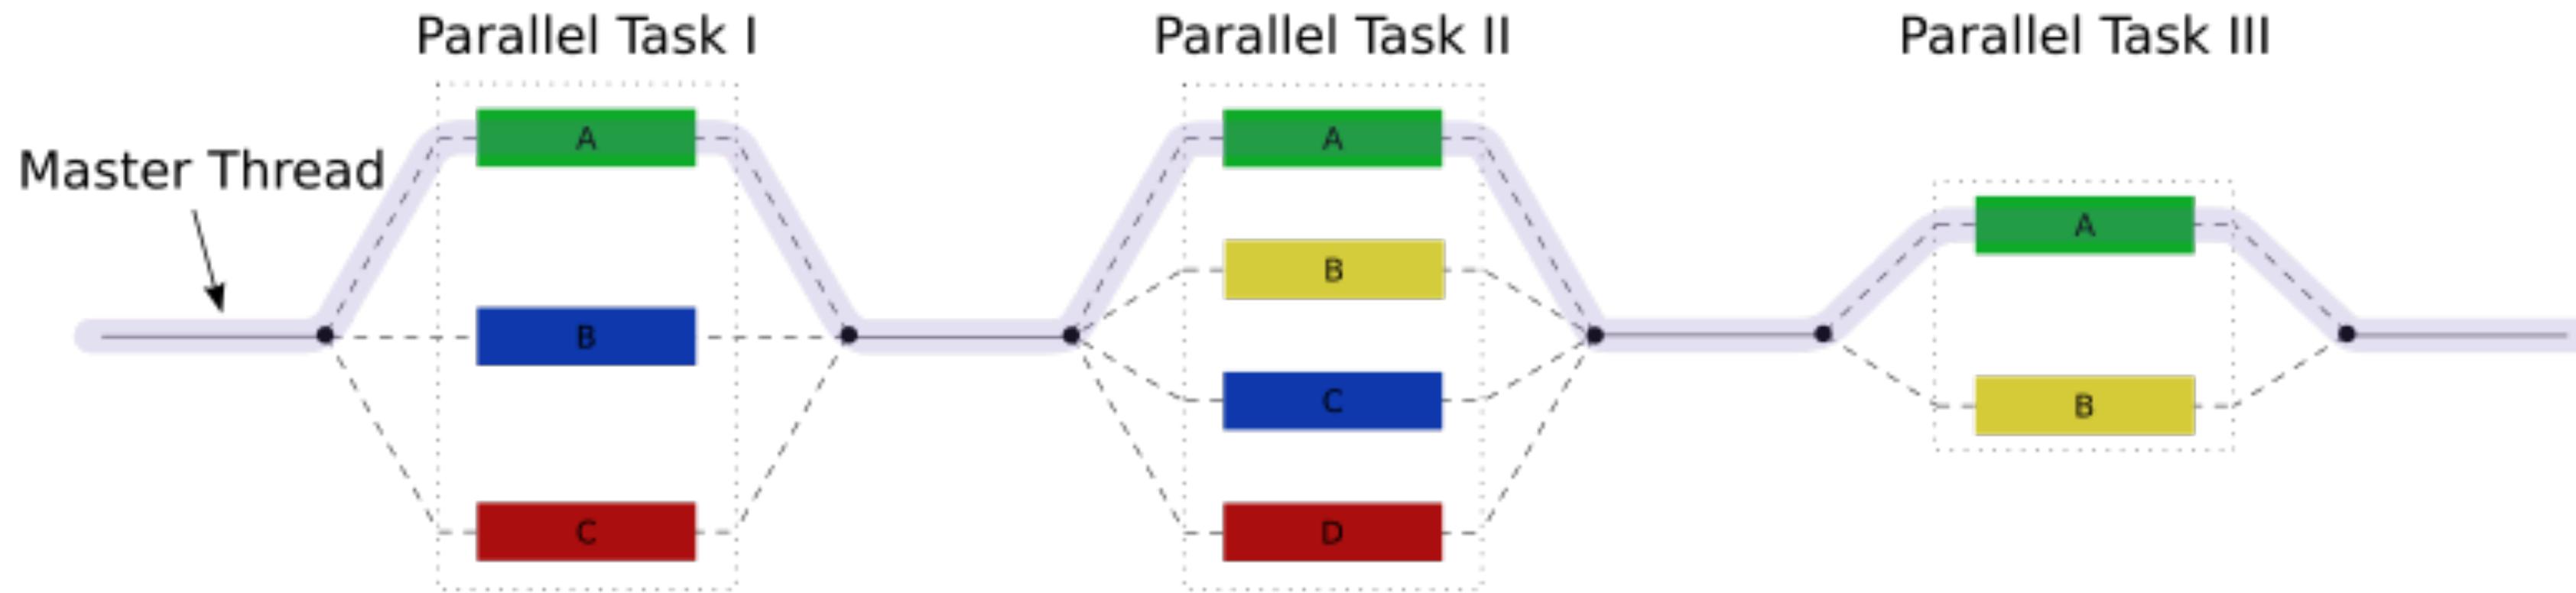
```
// ...
```

```
int fib (int n) {  
    if (n <= G) fib_seq (n);  
    int f1, f2;  
    #pragma omp task  
    f1 = fib (n-1);  
    f2 = fib (n-2);  
    #pragma omp taskwait  
    return f1 + f2;  
}
```

Controlling number of threads

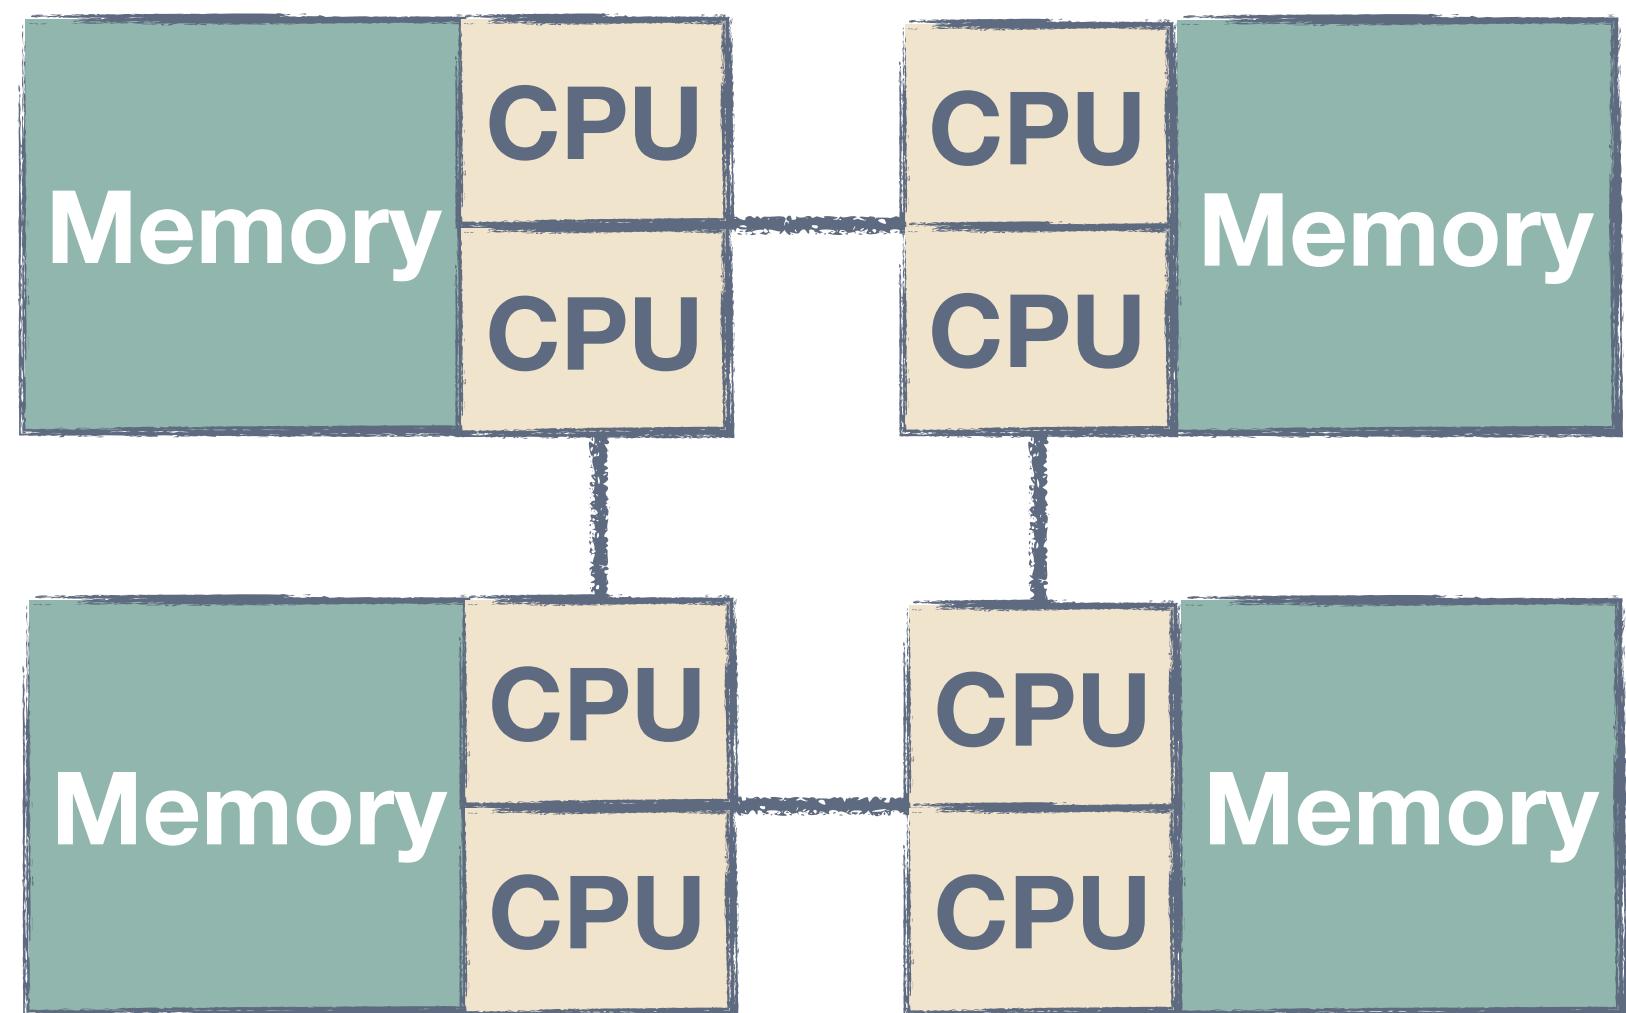


Controlling number of threads

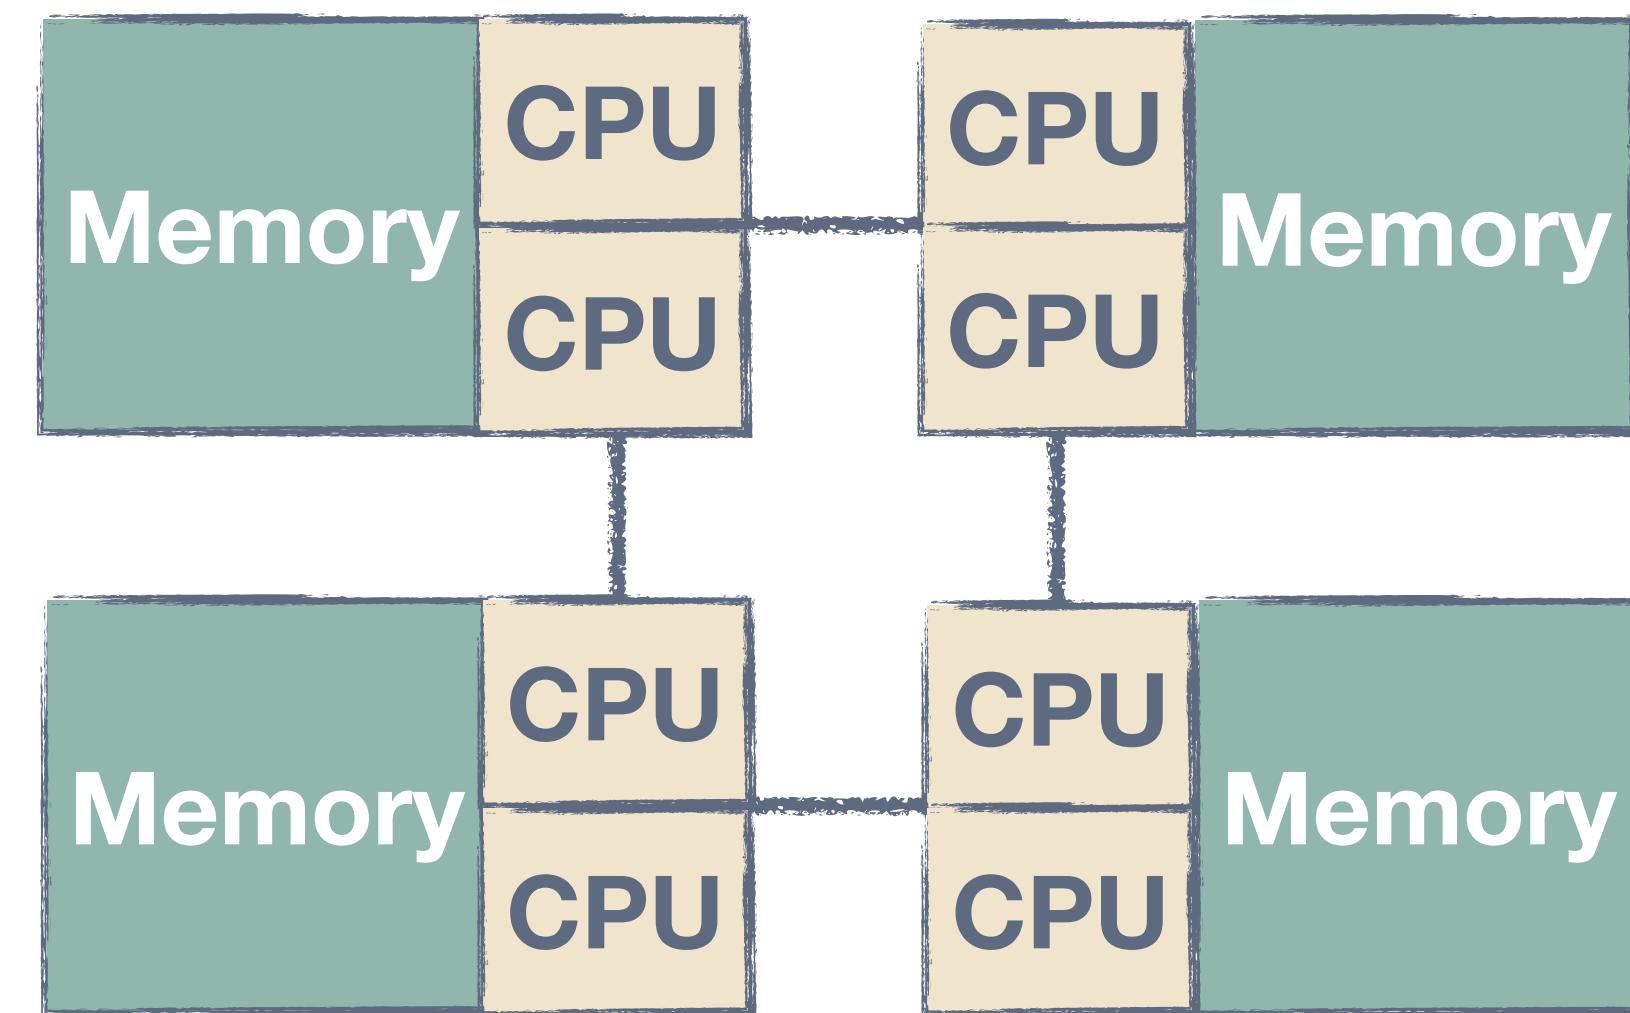


```
#pragma omp parallel num_threads(4)
{
    #pragma omp for shared (a,n) private (i)
    for (i = 0; i < n; ++i) {
        a[i] += i;
    } // Implicit barrier/join
} // Implicit barrier/join
```

NUMA: Linux “first-touch” policy



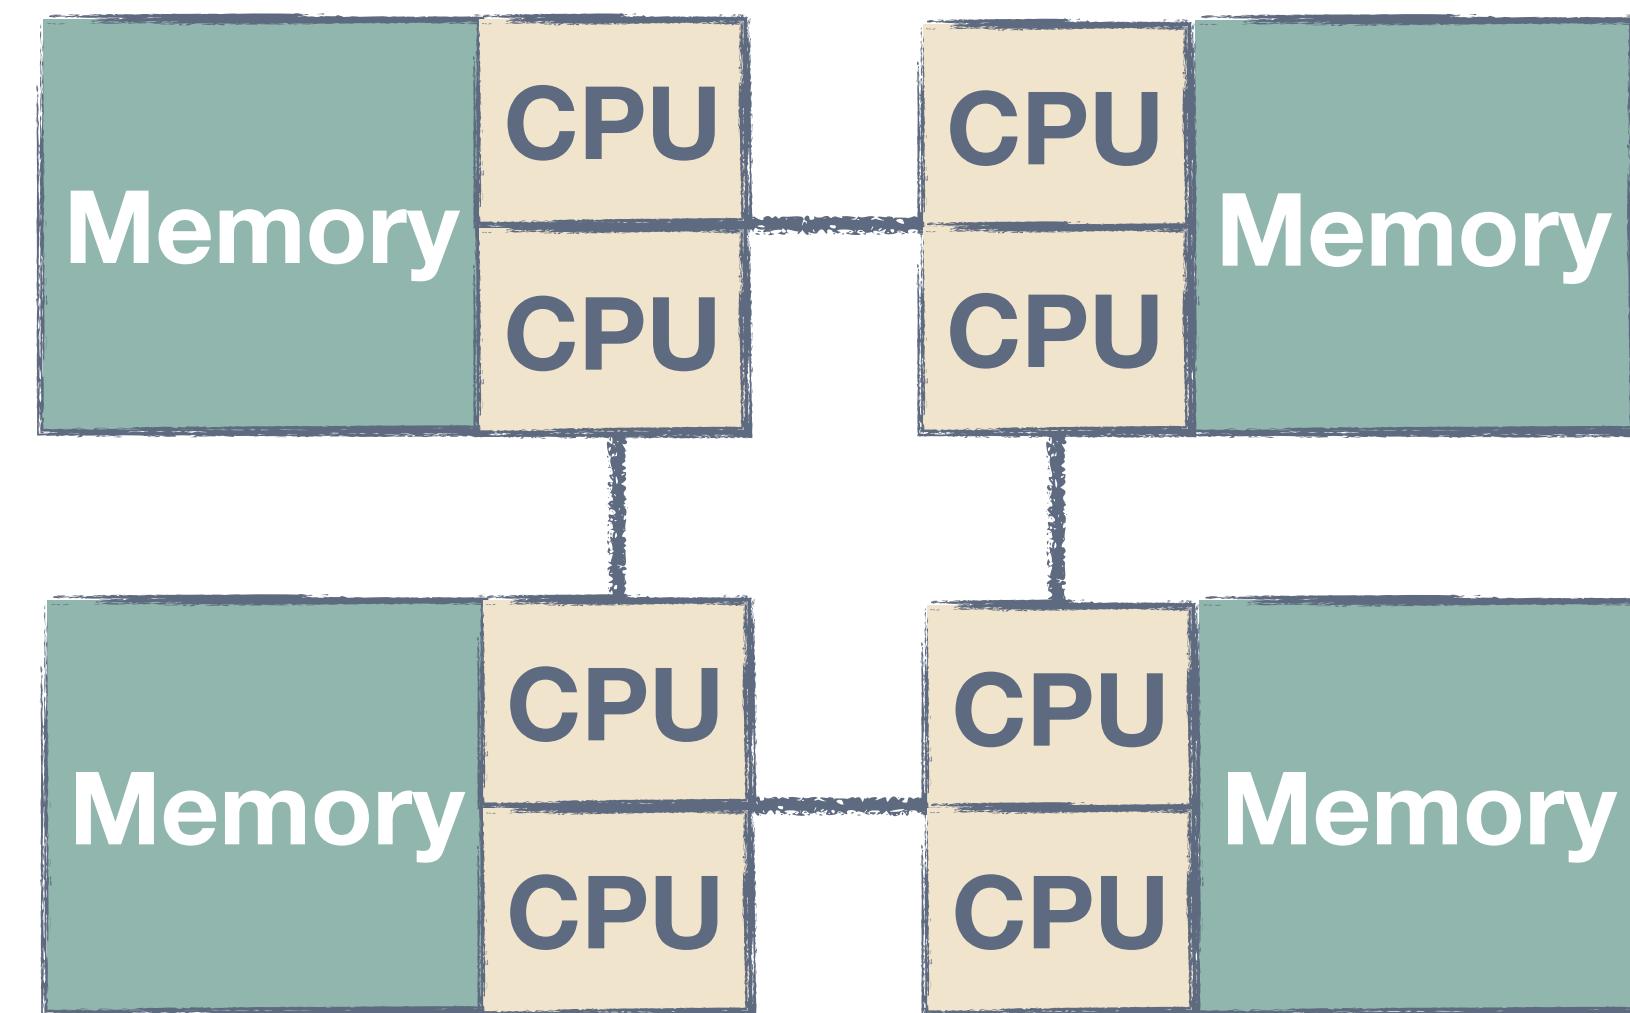
NUMA: Linux “first-touch” policy



```
#pragma omp parallel for shared (a,n) private (i)
for (i = 0; i < n; ++i) {
    a[i] += i;
} // Implicit barrier/join
```

NUMA: Linux “first-touch” policy

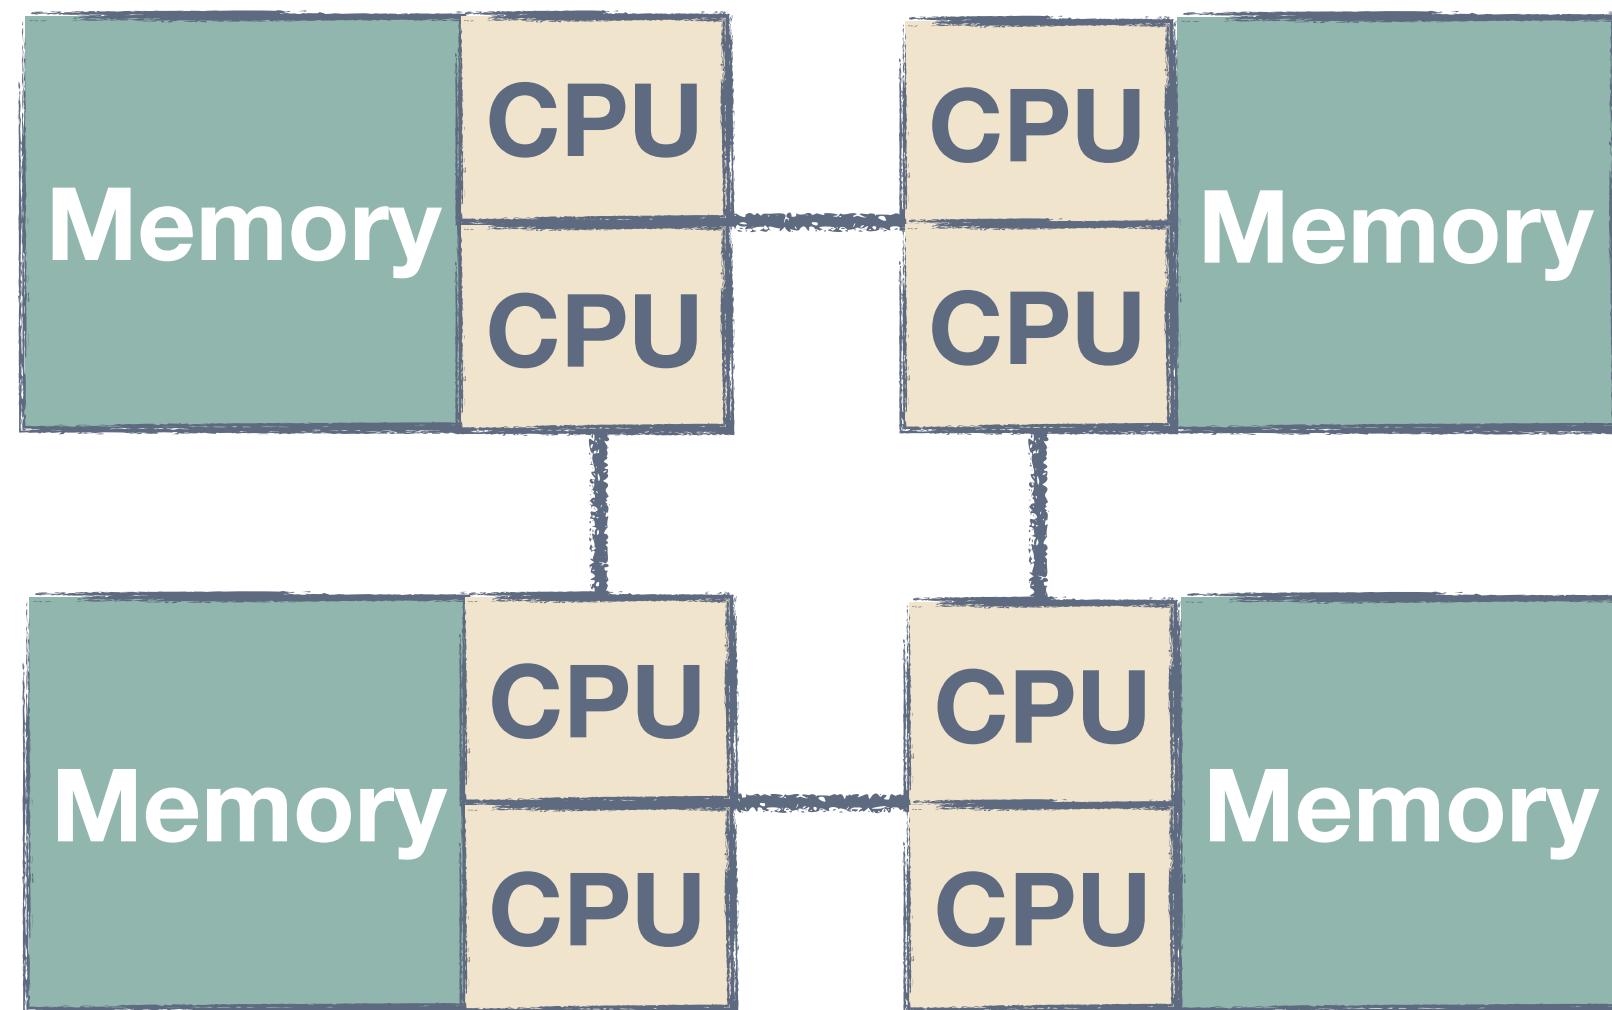
```
a = /* allocate memory */;  
  
for (i = 0; i < n; ++i) {  
    a[i] = /* assign initial value */;  
}
```



```
#pragma omp parallel for shared (a,n) private (i)  
for (i = 0; i < n; ++i) {  
    a[i] += i;  
} // Implicit barrier/join
```

NUMA: Linux “first-touch” policy

```
a = /* allocate memory */;  
#pragma omp parallel for schedule(static)  
for (i = 0; i < n; ++i) {  
    a[i] = /* assign initial value */;  
}
```



```
#pragma omp parallel for shared (a,n) private (i) schedule(static)  
for (i = 0; i < n; ++i) {  
    a[i] += i;  
} // Implicit barrier/join
```